

DonorsChoose

DonorsChoose.org receives hundreds of thousands of project proposals each year for classroom projects in need of funding. Right now, a large number of volunteers is needed to manually screen each submission before it's approved to be posted on the DonorsChoose.org website.

Next year, DonorsChoose.org expects to receive close to 500,000 project proposals. As a result, there are three main problems they need to solve:

- How to scale current manual processes and resources to screen 500,000 projects so that they can be posted as quickly and as efficiently as possible
- How to increase the consistency of project vetting across different volunteers to improve the experience for teachers
- How to focus volunteer time on the applications that need the most assistance

The goal of the competition is to predict whether or not a DonorsChoose.org project proposal submitted by a teacher will be approved, using the text of project descriptions as well as additional metadata about the project, teacher, and school. DonorsChoose.org can then use this information to identify projects most likely to need further review before approval.

About the DonorsChoose Data Set

The `train.csv` data set provided by DonorsChoose contains the following features:

Feature		
<code>project_id</code>		A unique identifier for the proposed project
<code>project_title</code>	<ul style="list-style-type: none">••	Title of the project Art Will
<code>project_grade_category</code>	<ul style="list-style-type: none">••••	Grade level of students for which the project is targeted

Feature	
	One or more (comma-separated) subject categories following enur
project_subject_categories	<ul style="list-style-type: none"> • • • • • • • •
	<ul style="list-style-type: none"> • •
school_state	State where school is located (Two- (https://en.wikipedia.org/wiki/List_of_U.S._state_abbreviations))
project_subject_subcategories	One or more (comma-separated) subject subcate <ul style="list-style-type: none"> • •
project_resource_summary	An explanation of the resources needed for t <ul style="list-style-type: none"> • My students need hands on literacy mar sen
project_essay_1	F
project_essay_2	Sec
project_essay_3	TI
project_essay_4	Fo
project_submitted_datetime	Datetime when project application was submitted. Ex:
teacher_id	A unique identifier for the teacher of the propos bdf8baa8fedef6b
	Teacher's title. One of the following <ul style="list-style-type: none"> • • • • • •
teacher_prefix	
teacher_number_of_previously_posted_projects	Number of project applications previously submitted

* See the section **Notes on the Essay Data** for more details about these features.

Additionally, the `resources.csv` data set provides more data about the resources required for each project. Each line in this file represents a resource required by a project:

Feature	Description
id	A <code>project_id</code> value from the <code>train.csv</code> file. Example: p036502
description	Description of the resource. Example: Tenor Saxophone Reeds, Box of 25
quantity	Quantity of the resource required. Example: 3
price	Price of the resource required. Example: 9.95

Note: Many projects require multiple resources. The `id` value corresponds to a `project_id` in `train.csv`, so you use it as a key to retrieve all resources needed for a project:

The data set contains the following label (the value you will attempt to predict):

Label	Description
<code>project_is_approved</code>	A binary flag indicating whether DonorsChoose approved the project. A value of <code>0</code> indicates the project was not approved, and a value of <code>1</code> indicates the project was approved.



Notes on the Essay Data

Prior to May 17, 2016, the prompts for the essays were as follows:

- `__project_essay_1:` "Introduce us to your classroom"
- `__project_essay_2:` "Tell us more about your students"
- `__project_essay_3:` "Describe how your students will use the materials you're requesting"
- `__project_essay_3:` "Close by sharing why your project will make a difference"

Starting on May 17, 2016, the number of essays was reduced from 4 to 2, and the prompts for the first 2 essays were changed to the following:

- `__project_essay_1:` "Describe your students: What makes your students special? Specific details about their background, your neighborhood, and your school are all helpful."
- `__project_essay_2:` "About your project: How will these materials make a difference in your students' learning and improve their school lives?"

For all projects with `project_submitted_datetime` of 2016-05-17 and later, the values of `project_essay_3` and `project_essay_4` will be NaN.

```
In [1]: %matplotlib inline
import warnings
warnings.filterwarnings("ignore")

import sqlite3
import pandas as pd
import numpy as np
import nltk
import string
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.feature_extraction.text import TfidfTransformer
from sklearn.feature_extraction.text import TfidfVectorizer

from sklearn.feature_extraction.text import CountVectorizer
from sklearn.metrics import confusion_matrix
from sklearn import metrics
from sklearn.metrics import roc_curve, auc
from nltk.stem.porter import PorterStemmer

import re
# Tutorial about Python regular expressions: https://pymotw.com/2/re/
import string
from nltk.corpus import stopwords
from nltk.stem import PorterStemmer
from nltk.stem.wordnet import WordNetLemmatizer

from gensim.models import Word2Vec
from gensim.models import KeyedVectors
import pickle

from tqdm import tqdm
import os

from plotly import plotly
import plotly.offline as offline
import plotly.graph_objs as go
offline.init_notebook_mode()
from collections import Counter
```

1.1 Reading Data

```
In [2]: project_data_ = pd.read_csv("train_new_data.csv")
resource_data_ = pd.read_csv("resources.csv")
```

```
In [3]: project_data=project_data_.head(25000)
resource_data=resource_data_.head(25000)
```

```
In [4]: print("Number of data points in train data", project_data.shape)
print('-'*50)
print("The attributes of data :", project_data.columns.values)
```

Number of data points in train data (25000, 17)

The attributes of data : ['Unnamed: 0' 'id' 'teacher_id' 'teacher_prefix' 'school_state'
 'project_submitted_datetime' 'project_grade_category'
 'project_subject_categories' 'project_subject_subcategories'
 'project_title' 'project_essay_1' 'project_essay_2' 'project_essay_3'
 'project_essay_4' 'project_resource_summary'
 'teacher_number_of_previously_posted_projects' 'project_is_approved']

```
In [5]: print("Number of data points in train data", resource_data.shape)
print(resource_data.columns.values)
resource_data.head(2)
```

Number of data points in train data (25000, 4)

['id' 'description' 'quantity' 'price']

```
Out[5]:
```

	id	description	quantity	price
0	p233245	LC652 - Lakeshore Double-Space Mobile Drying Rack	1	149.00
1	p069063	Bouncy Bands for Desks (Blue support pipes)	3	14.95

1.2 preprocessing of project_subject_categories

```

In [6]: categories = list(project_data['project_subject_categories'].values)
# remove special characters from list of strings python: https://stackoverflow.co
# https://www.geeksforgeeks.org/removing-stop-words-nltk-python/
# https://stackoverflow.com/questions/23669024/how-to-strip-a-specific-word-from
# https://stackoverflow.com/questions/8270092/remove-all-whitespace-in-a-string-
cat_list = []
for i in categories:
    temp = ""
    # consider we have text like this "Math & Science, Warmth, Care & Hunger"
    for j in i.split(','): # it will split it in three parts ["Math & Science",
        if 'The' in j.split(): # this will split each of the category based on sp
            j=j.replace('The','') # if we have the words "The" we are going to re
        j = j.replace(' ', '') # we are placeing all the ' '(space) with ''(empty,
        temp+=j.strip()+" " #" abc ".strip() will return "abc", remove the trail
        temp = temp.replace('&','_') # we are replacing the & value into
    cat_list.append(temp.strip())

project_data['clean_categories'] = cat_list
project_data.drop(['project_subject_categories'], axis=1, inplace=True)

from collections import Counter
my_counter = Counter()
for word in project_data['clean_categories'].values:
    my_counter.update(word.split())

cat_dict = dict(my_counter)
sorted_cat_dict = dict(sorted(cat_dict.items(), key=lambda kv: kv[1]))

```

1.3 preprocessing of project_subject_subcategories

```

In [7]: sub_categories = list(project_data['project_subject_subcategories'].values)
# remove special characters from list of strings python: https://stackoverflow.co

# https://www.geeksforgeeks.org/removing-stop-words-nltk-python/
# https://stackoverflow.com/questions/23669024/how-to-strip-a-specific-word-from
# https://stackoverflow.com/questions/8270092/remove-all-whitespace-in-a-string-
sub_cat_list = []
for i in sub_categories:
    temp = ""
    # consider we have text like this "Math & Science, Warmth, Care & Hunger"
    for j in i.split(','): # it will split it in three parts ["Math & Science",
        if 'The' in j.split(): # this will split each of the category based on sp
            j=j.replace('The','') # if we have the words "The" we are going to re
        j = j.replace(' ', '') # we are placing all the ' '(space) with ''(empty,
        temp +=j.strip()+" #" abc ".strip() will return "abc", remove the trail
        temp = temp.replace('&','_')
    sub_cat_list.append(temp.strip())

project_data['clean_subcategories'] = sub_cat_list
project_data.drop(['project_subject_subcategories'], axis=1, inplace=True)

# count of all the words in corpus python: https://stackoverflow.com/a/22898595/4
my_counter = Counter()
for word in project_data['clean_subcategories'].values:
    my_counter.update(word.split())

sub_cat_dict = dict(my_counter)
sorted_sub_cat_dict = dict(sorted(sub_cat_dict.items(), key=lambda kv: kv[1]))

```

preprocessing school state

```

In [8]: from collections import Counter
my_counter = Counter()
for word in project_data['school_state'].values:
    my_counter.update(word.split())
state_dict = dict(my_counter)
sorted_state_dict = dict(sorted(state_dict.items(), key=lambda kv: kv[1]))

```

preprocessing teacher prefix

```

In [9]: from collections import Counter
my_counter = Counter()
for word in project_data['teacher_prefix'].values:
    my_counter.update(word.split())
prefix_dict = dict(my_counter)
sorted_prefix_dict = dict(sorted(prefix_dict.items(), key=lambda kv: kv[1]))

```

preprocessing project grade category

```

In [10]: categories = list(project_data['project_grade_category'].values)
# remove special characters from list of strings python: https://stackoverflow.co

# https://www.geeksforgeeks.org/removing-stop-words-nltk-python/
# https://stackoverflow.com/questions/23669024/how-to-strip-a-specific-word-from
# https://stackoverflow.com/questions/8270092/remove-all-whitespace-in-a-string-
pgc_list = []
for i in categories:
    temp = ""
    # consider we have text like this "Math & Science, Warmth, Care & Hunger"
    for j in i.split(','): # it will split it in three parts ["Math & Science",
        if 'The' in j.split(): # this will split each of the category based on sp
            j=j.replace('The','') # if we have the words "The" we are going to re
        j = j.replace(' ', '') # we are placing all the ' '(space) with ''(empty)
        temp+=j.strip()+" " # " abc ".strip() will return "abc", remove the trail
        temp = temp.replace('&','_') # we are replacing the & value into
    pgc_list.append(temp.strip())

project_data['clean_pgc'] = pgc_list
project_data.drop(['project_grade_category'], axis=1, inplace=True)

from collections import Counter
my_counter = Counter()
for word in project_data['clean_pgc'].values:
    my_counter.update(word.split())

pgc_dict = dict(my_counter)
sorted_pgc_dict = dict(sorted(pgc_dict.items(), key=lambda kv: kv[1]))

```

```

In [11]: project_data.head(5)

```

```

Out[11]:

```

	Unnamed: 0	id	teacher_id	teacher_prefix	school_state	project_su
0	0	p036502	484aaf11257089a66cfedc9461c6bd0a	Ms.	NV	
1	3	p185307	525fdbb6ec7f538a48beebaa0a51b24f	Mr.	NC	
2	4	p013780	a63b5547a7239eae4c1872670848e61a	Mr.	CA	
3	5	p063374	403c6783e9286e51ab318fba40f8d729	Mrs.	DE	
4	6	p103285	4e156c5fb3eea2531601c8736f3751a7	Mrs.	MO	

1.3 Text preprocessing

```
In [12]: # merge two column text dataframe:
project_data["essay"] = project_data["project_essay_1"].map(str) + \
    project_data["project_essay_2"].map(str) + \
    project_data["project_essay_3"].map(str) + \
    project_data["project_essay_4"].map(str)
```

```
In [13]: project_data.head(2)
```

```
Out[13]:
```

	Unnamed: 0	id	teacher_id	teacher_prefix	school_state	project_sub
--	------------	----	------------	----------------	--------------	-------------

0	0	p036502	484aaf11257089a66cfedc9461c6bd0a	Ms.	NV	
1	3	p185307	525fdbb6ec7f538a48beebaa0a51b24f	Mr.	NC	

Decontracting function for sentence

```
In [14]: # https://stackoverflow.com/a/47091490/4084039
import re

def decontracted(phrase):
    # specific
    phrase = re.sub(r"won't", "will not", phrase)
    phrase = re.sub(r"can't", "can not", phrase)

    # general
    phrase = re.sub(r"n't", " not", phrase)
    phrase = re.sub(r"\ 're", " are", phrase)
    phrase = re.sub(r"\ 's", " is", phrase)
    phrase = re.sub(r"\ 'd", " would", phrase)
    phrase = re.sub(r"\ 'll", " will", phrase)
    phrase = re.sub(r"\ 't", " not", phrase)
    phrase = re.sub(r"\ 've", " have", phrase)
    phrase = re.sub(r"\ 'm", " am", phrase)
    return phrase
```

```
In [15]: # https://gist.github.com/sebleier/554280
# we are removing the words from the stop words list: 'no', 'nor', 'not'
stopwords= ['i', 'me', 'my', 'myself', 'we', 'our', 'ours', 'ourselves', 'you', 'you\'ll', 'you\'d', 'your', 'yours', 'yourself', 'yourselves', 'he', 'she', 'she\'s', 'her', 'hers', 'herself', 'it', 'it\'s', 'its', 'itself', 'theirs', 'themselves', 'what', 'which', 'who', 'whom', 'this', 'that', 'am', 'is', 'are', 'was', 'were', 'be', 'been', 'being', 'have', 'has', 'did', 'doing', 'a', 'an', 'the', 'and', 'but', 'if', 'or', 'because', 'at', 'by', 'for', 'with', 'about', 'against', 'between', 'into', 'through', 'above', 'below', 'to', 'from', 'up', 'down', 'in', 'out', 'on', 'off', 'then', 'once', 'here', 'there', 'when', 'where', 'why', 'how', 'all', 'most', 'other', 'some', 'such', 'only', 'own', 'same', 'so', 'than', 's', 't', 'can', 'will', 'just', 'don', 'don\'t', 'should', 'should\'ve', 've', 'y', 'ain', 'aren', 'aren\'t', 'couldn', 'couldn\'t', 'didn', 'didn\'t', 'hadn\'t', 'hasn', 'hasn\'t', 'haven', 'haven\'t', 'isn', 'isn\'t', 'ma', 'mustn\'t', 'needn', 'needn\'t', 'shan', 'shan\'t', 'shouldn', 'shouldn\'t', 'won', 'won\'t', 'wouldn', 'wouldn\'t']
```

```
In [16]: # Combining all the above students
from tqdm import tqdm
preprocessed_essays = []
# tqdm is for printing the status bar
for sentence in tqdm(project_data['essay'].values):
    sent = decontracted(sentence)
    sent = sent.replace('\\r', ' ')
    sent = sent.replace('\\\"', ' ')
    sent = sent.replace('\\n', ' ')
    sent = re.sub('[^A-Za-z0-9]+', ' ', sent)
    sent=sent.lower()
    # https://gist.github.com/sebleier/554280
    sent = ' '.join(e for e in sent.split() if e.lower() not in stopwords)
    preprocessed_essays.append(sent.lower().strip())
```

100%|██████████| 25000/25000 [00:23<00:00, 1050.52it/s]

```
In [17]: # after preprocessing
preprocessed_essays[2000]
```

```
Out[17]: 'bilingual first grade students full joy eager learn classroom place daily growth constant challenge discovery students spend year learning foundations reading writing math order succeed lives quickly becoming independent learners taking information learned apply multiple activities allow use imagination high level thinking skills teacher low income high poverty school district students faced several challenges classroom personal folders used every day reading writing math classes provide students personal space using folders help students focus work not neighbor students able use dividers whole group independent small group time instruction generous donation project improve students self confidence independence donating project not help improve increase student attention focus ultimately help increase academic achievementnannan'
```

```
In [18]: project_data["clean_essays"] = preprocessed_essays

project_data.drop(['essay'], axis=1, inplace=True)
```

1.4 Preprocessing of `project_title`

```
In [19]: preprocessed_pt = []
for titles in tqdm(project_data["project_title"]):
    title = decontracted(titles)
    title = title.replace('\r', ' ')
    title = title.replace('\n', ' ')
    title = title.replace('\n', ' ')
    title = re.sub('[^A-Za-z0-9]+', ' ', title)
    title = ' '.join(f for f in title.split() if f not in stopwords)
    preprocessed_pt.append(title.lower().strip())
```

100%|██████████| 25000/25000 [00:01<00:00, 20236.37it/s]

```
In [20]: project_data["clean_pt"] = preprocessed_pt
project_data.drop(['project_title'], axis=1, inplace=True)
```

number of words in title

```
In [21]: title_word_count = []
for i in project_data["clean_pt"] :
    j = len(i.split())
    title_word_count.append(j)
project_data["title_word_count"] = title_word_count
project_data.head(5)
```

```
Out[21]:
```

	Unnamed: 0	id	teacher_id	teacher_prefix	school_state	project_su
0	0	p036502	484aaf11257089a66cfedc9461c6bd0a	Ms.	NV	
1	3	p185307	525fdbb6ec7f538a48beebaa0a51b24f	Mr.	NC	
2	4	p013780	a63b5547a7239eae4c1872670848e61a	Mr.	CA	
3	5	p063374	403c6783e9286e51ab318fba0f8d729	Mrs.	DE	
4	6	p103285	4e156c5fb3eea2531601c8736f3751a7	Mrs.	MO	

number of words in essay

```
In [22]: essay_word_count = []
for i in project_data["clean_essays"] :
    j = len(i.split())
    essay_word_count.append(j)
project_data["essay_word_count"] = essay_word_count
project_data.head(5)
```

```
Out[22]:
```

	Unnamed: 0	id	teacher_id	teacher_prefix	school_state	project
0	0	p036502	484aaf11257089a66cfedc9461c6bd0a	Ms.	NV	
1	3	p185307	525fdbb6ec7f538a48beebaa0a51b24f	Mr.	NC	
2	4	p013780	a63b5547a7239eae4c1872670848e61a	Mr.	CA	

Calculate Sentiment Scores for the essays

```
In [23]: import nltk
from nltk.sentiment.vader import SentimentIntensityAnalyzer

analyser = SentimentIntensityAnalyzer()
```

```
In [24]: neg = []
pos = []
neu = []
compound = []
for i in tqdm(project_data["clean_essays"]):
    j = analyser.polarity_scores(i)['neg']
    k = analyser.polarity_scores(i)['pos']
    l = analyser.polarity_scores(i)['neu']
    m = analyser.polarity_scores(i)['compound']
    neg.append(j)
    pos.append(k)
    neu.append(l)
    compound.append(m)
```

100%|██████████| 25000/25000 [04:52<00:00, 85.40it/s]

```
In [25]: project_data["neg"] = neg
project_data["pos"] = pos
project_data["neu"] = neu
project_data["compound"] = compound
```

```
In [26]: project_data.head(2)
```

```
Out[26]:
```

	Unnamed: 0	id	teacher_id	teacher_prefix	school_state	project_sub
0	0	p036502	484aaf11257089a66cfedc9461c6bd0a	Ms.	NV	
1	3	p185307	525fdbb6ec7f538a48beebaa0a51b24f	Mr.	NC	

2 rows × 24 columns

```
In [27]: project_data=project_data.head(25000)
```

```
In [28]: project_data.shape
```

```
Out[28]: (25000, 24)
```

Splitting data as train ,test and CV

```
In [29]: from sklearn.model_selection import train_test_split
S_train, S_test, y_train, y_test = train_test_split(project_data,
project_data['project_is_approved'], test_size=0.33, stratify = project_data['pro
])
S_train, S_cv, y_train, y_cv = train_test_split(S_train, y_train, test_size=0.30,
```

```
In [30]: S_train.drop(['project_is_approved'], axis=1, inplace=True)
S_test.drop(['project_is_approved'], axis=1, inplace=True)
S_cv.drop(['project_is_approved'], axis=1, inplace=True)
```

```
In [31]: print(S_train.shape)
print(S_test.shape)
print(S_cv.shape)
```

```
(11725, 23)
(8250, 23)
(5025, 23)
```

1.5 Preparing data for models

```
In [32]: project_data.columns
```

```
Out[32]: Index(['Unnamed: 0', 'id', 'teacher_id', 'teacher_prefix', 'school_state',  
              'project_submitted_datetime', 'project_essay_1', 'project_essay_2',  
              'project_essay_3', 'project_essay_4', 'project_resource_summary',  
              'teacher_number_of_previously_posted_projects', 'project_is_approved',  
              'clean_categories', 'clean_subcategories', 'clean_pgc', 'clean_essays',  
              'clean_pt', 'title_word_count', 'essay_word_count', 'neg', 'pos', 'neu',  
              'compound'],  
              dtype='object')
```

we are going to consider

- school_state : categorical data
- clean_categories : categorical data
- clean_subcategories : categorical data
- project_grade_category : categorical data
- teacher_prefix : categorical data

- project_title : text data
- text : text data
- project_resource_summary: text data (optinal)

- quantity : numerical (optinal)
- teacher_number_of_previously_posted_projects : numerical
- price : numerical

1.5.1 Vectorizing Categorical data

- <https://www.appliedaicourse.com/course/applied-ai-course-online/lessons/handling-categorical-and-numerical-features/> (<https://www.appliedaicourse.com/course/applied-ai-course-online/lessons/handling-categorical-and-numerical-features/>)

RESPONSE CODING FOR FEATURIZATION

```
In [33]: def response_code_dict(alpha, feature, df):
count=S_train[feature].value_counts()
resp_code=dict()
for i,deno in count.items():
    lis=[]
    for j in range(0,2):
        class_count=S_train.loc[(y_train==j) & (S_train[feature]==i) ]
        lis.append((class_count.shape[0]+10*alpha)/deno+(90*alpha))
    resp_code[i]=lis
return resp_code

def response_code_value(alpha, feature, df):

    get_values_dict = response_code_dict(alpha, feature, df)
    value_count = S_train[feature].value_counts()
    get_values_features = []
    for index, row in df.iterrows():
        if row[feature] in dict(value_count).keys():
            get_values_features.append(get_values_dict[row[feature]])
        else:
            get_values_features.append([1/2,1/2])
    return get_values_features
```

VECTORIZING CLEAN CATEGORIES USING RESPONSE CODING

```
In [34]: #response-coding of the Gene feature
# alpha is used for Laplace smoothing
alpha = 1
# train gene feature
train_clean_cat_feature_responseCoding = np.array(response_code_value(alpha, "clean", train_data))
# test gene feature
test_clean_cat_feature_responseCoding = np.array(response_code_value(alpha, "clean", test_data))
# cross validation gene feature
cv_clean_cat_feature_responseCoding = np.array(response_code_value(alpha, "clean", cv_data))
```

```
In [35]: print(train_clean_cat_feature_responseCoding.shape)
print(test_clean_cat_feature_responseCoding.shape)
print(cv_clean_cat_feature_responseCoding.shape)
```

```
(11725, 2)
(8250, 2)
(5025, 2)
```

```
In [36]: from scipy import sparse
```

```
In [37]: e1=sparse.csr_matrix(train_clean_cat_feature_responseCoding.shape)
e2=sparse.csr_matrix(test_clean_cat_feature_responseCoding.shape)
e3=sparse.csr_matrix(cv_clean_cat_feature_responseCoding.shape)
```

VECTORIZING CLEAN SUBCATEGORIES USING RESPONSE CODING

```
In [38]: train_clean_subcat_feature_responseCoding = np.array(response_code_value(alpha, "train_clean_subcat_feature_responseCoding"))
# test gene feature
test_clean_subcat_feature_responseCoding = np.array(response_code_value(alpha, "test_clean_subcat_feature_responseCoding"))
# cross validation gene feature
cv_clean_subcat_feature_responseCoding = np.array(response_code_value(alpha, "cv_clean_subcat_feature_responseCoding"))
```

```
In [39]: print(train_clean_subcat_feature_responseCoding.shape)
print(test_clean_subcat_feature_responseCoding.shape)
print(cv_clean_subcat_feature_responseCoding.shape)

(11725, 2)
(8250, 2)
(5025, 2)
```

```
In [40]: d1=sparse.csr_matrix(train_clean_subcat_feature_responseCoding)
d2=sparse.csr_matrix(test_clean_subcat_feature_responseCoding)
d3=sparse.csr_matrix(cv_clean_subcat_feature_responseCoding)
```

VECTORIZING SCHOOL STATE USING RESPONSE CODING

```
In [41]: # you can do the similar thing with state, teacher_prefix and project_grade_category
train_state_feature_responseCoding = np.array(response_code_value(alpha, "train_state_feature_responseCoding"))
# test gene feature
test_state_feature_responseCoding = np.array(response_code_value(alpha, "test_state_feature_responseCoding"))
# cross validation gene feature
cv_state_feature_responseCoding = np.array(response_code_value(alpha, "cv_state_feature_responseCoding"))
```

```
In [42]: print(train_state_feature_responseCoding.shape)
print(test_state_feature_responseCoding.shape)
print(cv_state_feature_responseCoding.shape)

(11725, 2)
(8250, 2)
(5025, 2)
```

```
In [43]: c1=sparse.csr_matrix(train_state_feature_responseCoding)
c2=sparse.csr_matrix(test_state_feature_responseCoding)
c3=sparse.csr_matrix(cv_state_feature_responseCoding)
```

VECTORIZING TEACHER PREFIX USING RESPONSE CODING

```
In [44]: train_prefix_feature_responseCoding = np.array(response_code_value(alpha, "train_prefix_feature_responseCoding"))
# test gene feature
test_prefix_feature_responseCoding = np.array(response_code_value(alpha, "test_prefix_feature_responseCoding"))
# cross validation gene feature
cv_prefix_feature_responseCoding = np.array(response_code_value(alpha, "cv_prefix_feature_responseCoding"))
```



```
In [45]: print(train_prefix_feature_responseCoding.shape)
print(test_prefix_feature_responseCoding.shape)
print(cv_prefix_feature_responseCoding.shape)

(11725, 2)
(8250, 2)
(5025, 2)
```

```
In [46]: b1=sparse.csr_matrix(train_prefix_feature_responseCoding)
b2=sparse.csr_matrix(test_prefix_feature_responseCoding)
b3=sparse.csr_matrix(cv_prefix_feature_responseCoding)
```

VECTORIZING PROJECT GRADE CATEGORY USING RESPONSE CODING

```
In [47]: train_clean_pgc_feature_responseCoding = np.array(response_code_value(alpha, "clean_essays"))
# test gene feature
test_clean_pgc_feature_responseCoding = np.array(response_code_value(alpha, "clean_essays"))
# cross validation gene feature
cv_clean_pgc_feature_responseCoding = np.array(response_code_value(alpha, "clean_essays"))
```

```
In [48]: print(train_clean_pgc_feature_responseCoding.shape)
print(test_clean_pgc_feature_responseCoding.shape)
print(cv_clean_pgc_feature_responseCoding.shape)

(11725, 2)
(8250, 2)
(5025, 2)
```

```
In [49]: from scipy import sparse
a1=sparse.csr_matrix(train_clean_pgc_feature_responseCoding)
a2=sparse.csr_matrix(test_clean_pgc_feature_responseCoding)
a3=sparse.csr_matrix(cv_clean_pgc_feature_responseCoding)
```

1.5.2 Vectorizing Text data

1.5.2.1 Bag of words

```
In [50]: # We are considering only the words which appeared in at least 10 documents(rows)
vectorizer_bow = CountVectorizer()
text_bow = vectorizer_bow.fit_transform(S_train["clean_essays"])
print("Shape of matrix after one hot encoding ",text_bow.shape)
```

Shape of matrix after one hot encoding (11725, 23885)

```
In [51]: text_bow_test = vectorizer_bow.transform(S_test["clean_essays"])
print("Shape of matrix after one hot encoding ",text_bow_test.shape)
```

Shape of matrix after one hot encoding (8250, 23885)

```
In [52]: text_bow_cv = vectorizer_bow.transform(S_cv["clean_essays"])
print("Shape of matrix after one hot encoding ",text_bow_cv.shape)
```

Shape of matrix after one hot encoding (5025, 23885)

```
In [53]: vectorizer_title_bow = CountVectorizer()
title_bow_train= vectorizer_title_bow.fit_transform(S_train["clean_pt"])
print("Shape of matrix after one hot encoding ",title_bow_train.shape)
```

Shape of matrix after one hot encoding (11725, 5834)

```
In [54]: title_bow_test = vectorizer_title_bow.transform(S_test["clean_pt"])
print("Shape of matrix after one hot encoding ",title_bow_test.shape)
```

Shape of matrix after one hot encoding (8250, 5834)

```
In [55]: title_bow_cv = vectorizer_title_bow.transform(S_cv["clean_pt"])
print("Shape of matrix after one hot encoding ",title_bow_cv.shape)
```

Shape of matrix after one hot encoding (5025, 5834)

1.5.2.2 TFIDF vectorizer

```
In [56]: from sklearn.feature_extraction.text import TfidfVectorizer
vectorizer_tfidf_essay = TfidfVectorizer()
vectorizer_tfidf_essay.fit(S_train["clean_essays"])
text_tfidf_train = vectorizer_tfidf_essay.transform(S_train["clean_essays"])
print("Shape of matrix after one hot encoding ",text_tfidf_train.shape)
```

Shape of matrix after one hot encoding (11725, 23885)

```
In [57]: text_tfidf_test = vectorizer_tfidf_essay.transform(S_test["clean_essays"])
print("Shape of matrix after one hot encoding ",text_tfidf_test.shape)
```

Shape of matrix after one hot encoding (8250, 23885)

```
In [58]: text_tfidf_cv = vectorizer_tfidf_essay.transform(S_cv["clean_essays"])
print("Shape of matrix after one hot encoding ",text_tfidf_cv.shape)
```

Shape of matrix after one hot encoding (5025, 23885)

```
In [59]: vectorizer_tfidf_title = TfidfVectorizer()
vectorizer_tfidf_title.fit(S_train["clean_pt"])
title_tfidf_train = vectorizer_tfidf_title.transform(S_train["clean_pt"])
print("Shape of matrix after one hot encoding ",title_tfidf_train.shape)
```

Shape of matrix after one hot encoding (11725, 5834)

```
In [60]: title_tfidf_test = vectorizer_tfidf_title.transform(S_test["clean_pt"])
print("Shape of matrix after one hot encoding ",title_tfidf_test.shape)
```

Shape of matrix after one hot encoding (8250, 5834)

```
In [61]: title_tfidf_cv = vectorizer_tfidf_title.transform(S_cv["clean_pt"])  
         print("Shape of matrix after one hot encoding ",title_tfidf_cv.shape)
```

Shape of matrix after one hot encoding (5025, 5834)

1.5.2.3 Using Pretrained Models: Avg W2V

In [62]:

```

# Reading glove vectors in python: https://stackoverflow.com/a/38230349/4084039
def loadGloveModel(gloveFile):
    print ("Loading Glove Model")
    f = open(gloveFile, 'r', encoding="utf8")
    model = {}
    for line in tqdm(f):
        splitLine = line.split()
        word = splitLine[0]
        embedding = np.array([float(val) for val in splitLine[1:]])
        model[word] = embedding
    print ("Done.", len(model), " words loaded!")
    return model
model = loadGloveModel('glove.42B.300d.txt')

words = []
for i in preprocessed_essays:
    words.extend(i.split(' '))

for i in preprocessed_pt:
    words.extend(i.split(' '))
print("all the words in the coupus", len(words))
words = set(words)
print("the unique words in the coupus", len(words))

inter_words = set(model.keys()).intersection(words)
print("The number of words that are present in both glove vectors and our coupus",
      len(inter_words), "(", np.round(len(inter_words)/len(words)*100, 3), "%)")

words_courpus = {}
words_glove = set(model.keys())
for i in words:
    if i in words_glove:
        words_courpus[i] = model[i]
print("word 2 vec length", len(words_courpus))

# stronging variables into pickle files python: http://www.jessicayung.com/how-to-
import pickle
with open('glove_vectors', 'wb') as f:
    pickle.dump(words_courpus, f)

```

Loading Glove Model

279727it [01:21, 3450.27it/s]

Done. 279727 words loaded!

all the words in the coupus 3566468

the unique words in the coupus 32982

The number of words that are present in both glove vectors and our coupus 287

48 (87.163 %)
word 2 vec length 28748

```
In [63]: # stronging variables into pickle files python: http://www.jessicayung.com/how-to
# make sure you have the glove_vectors file
with open('glove_vectors', 'rb') as f:
    model = pickle.load(f)
    glove_words = set(model.keys())
```

```
In [64]: # average Word2Vec
# compute average word2vec for each review.
avg_w2v_vectors_train = []; # the avg-w2v for each sentence/review is stored in t
for sentence in tqdm(S_train["clean_essays"]): # for each review/sentence
    vector = np.zeros(300) # as word vectors are of zero length
    cnt_words = 0; # num of words with a valid vector in the sentence/review
    for word in sentence.split(): # for each word in a review/sentence
        if word in glove_words:
            vector += model[word]
            cnt_words += 1
    if cnt_words != 0:
        vector /= cnt_words
    avg_w2v_vectors_train.append(vector)

print(len(avg_w2v_vectors_train))
print(len(avg_w2v_vectors_train[0]))
```

100%|██████████| 11725/11725 [00:13<00:00, 866.80it/s]

11725
300

```
In [65]: avg_w2v_vectors_test = []; # the avg-w2v for each sentence/review is stored in t
for sentence in tqdm(S_test["clean_essays"]): # for each review/sentence
    vector = np.zeros(300) # as word vectors are of zero length
    cnt_words = 0; # num of words with a valid vector in the sentence/review
    for word in sentence.split(): # for each word in a review/sentence
        if word in glove_words:
            vector += model[word]
            cnt_words += 1
    if cnt_words != 0:
        vector /= cnt_words
    avg_w2v_vectors_test.append(vector)

print(len(avg_w2v_vectors_test))
print(len(avg_w2v_vectors_test[0]))
```

100%|██████████| 8250/8250 [00:03<00:00, 2062.72it/s]

8250
300

```
In [66]: avg_w2v_vectors_cv = []; # the avg-w2v for each sentence/review is stored in this
for sentence in tqdm(S_cv["clean_essays"]): # for each review/sentence
    vector = np.zeros(300) # as word vectors are of zero length
    cnt_words = 0; # num of words with a valid vector in the sentence/review
    for word in sentence.split(): # for each word in a review/sentence
        if word in glove_words:
            vector += model[word]
            cnt_words += 1
    if cnt_words != 0:
        vector /= cnt_words
    avg_w2v_vectors_cv.append(vector)

print(len(avg_w2v_vectors_cv))
print(len(avg_w2v_vectors_cv[0]))
```

100%|██████████| 5025/5025 [00:03<00:00, 1627.48it/s]

5025

300

```
In [67]: avg_w2v_title_train = []; # the avg-w2v for each sentence/review is stored in this
for sentence in tqdm(S_train["clean_pt"]): # for each review/sentence
    vector = np.zeros(300) # as word vectors are of zero length
    cnt_words = 0; # num of words with a valid vector in the sentence/review
    for word in sentence.split(): # for each word in a review/sentence
        if word in glove_words:
            vector += model[word]
            cnt_words += 1
    if cnt_words != 0:
        vector /= cnt_words
    avg_w2v_title_train.append(vector)

print(len(avg_w2v_title_train))
print(len(avg_w2v_title_train[0]))
```

100%|██████████| 11725/11725 [00:00<00:00, 41469.88it/s]

11725

300

```
In [68]: avg_w2v_title_test = []; # the avg-w2v for each sentence/review is stored in this
for sentence in tqdm(S_test["clean_pt"]): # for each review/sentence
    vector = np.zeros(300) # as word vectors are of zero length
    cnt_words = 0; # num of words with a valid vector in the sentence/review
    for word in sentence.split(): # for each word in a review/sentence
        if word in glove_words:
            vector += model[word]
            cnt_words += 1
    if cnt_words != 0:
        vector /= cnt_words
    avg_w2v_title_test.append(vector)

print(len(avg_w2v_vectors_test))
print(len(avg_w2v_vectors_test[0]))
```

100%|██████████| 8250/8250 [00:00<00:00, 27750.87it/s]

8250

300

```
In [69]: avg_w2v_title_cv = []; # the avg-w2v for each sentence/review is stored in this
for sentence in tqdm(S_cv["clean_pt"]): # for each review/sentence
    vector = np.zeros(300) # as word vectors are of zero length
    cnt_words = 0; # num of words with a valid vector in the sentence/review
    for word in sentence.split(): # for each word in a review/sentence
        if word in glove_words:
            vector += model[word]
            cnt_words += 1
    if cnt_words != 0:
        vector /= cnt_words
    avg_w2v_title_cv.append(vector)

print(len(avg_w2v_title_cv))
print(len(avg_w2v_title_cv[0]))
```

100%|██████████| 5025/5025 [00:00<00:00, 32228.40it/s]

5025

300

1.5.2.3 Using Pretrained Models: TFIDF weighted W2V

```
In [70]: # S = ["abc def pqr", "def def def abc", "pqr pqr def"]
tfidf_model = TfidfVectorizer()
tfidf_model.fit(S_train["clean_essays"])
# we are converting a dictionary with word as a key, and the idf as a value
dictionary = dict(zip(tfidf_model.get_feature_names(), list(tfidf_model.idf_)))
tfidf_words = set(tfidf_model.get_feature_names())
```

```
In [71]: tfidf_w2v_vectors_train = []; # the avg-w2v for each sentence/review is stored in
for sentence in tqdm(S_train["clean_essays"]): # for each review/sentence
    vector = np.zeros(300) # as word vectors are of zero length
    tf_idf_weight = 0; # num of words with a valid vector in the sentence/review
    for word in sentence.split(): # for each word in a review/sentence
        if (word in glove_words) and (word in tfidf_words):
            vec = model[word] # getting the vector for each word
            # here we are multiplying idf value(dictionary[word]) and the tf value
            tf_idf = dictionary[word]*(sentence.count(word)/len(sentence.split()))
            vector += (vec * tf_idf) # calculating tfidf weighted w2v
            tf_idf_weight += tf_idf
    if tf_idf_weight != 0:
        vector /= tf_idf_weight
    tfidf_w2v_vectors_train.append(vector)

print(len(tfidf_w2v_vectors_train))
print(len(tfidf_w2v_vectors_train[0]))
```

100%|██████████| 11725/11725 [00:39<00:00, 293.87it/s]

11725

300

```
In [72]: tfidf_w2v_vectors_test= []; # the avg-w2v for each sentence/review is stored in
for sentence in tqdm(S_test["clean_essays"]): # for each review/sentence
    vector = np.zeros(300) # as word vectors are of zero length
    tf_idf_weight = 0; # num of words with a valid vector in the sentence/review
    for word in sentence.split(): # for each word in a review/sentence
        if (word in glove_words) and (word in tfidf_words):
            vec = model[word] # getting the vector for each word
            # here we are multiplying idf value(dictionary[word]) and the tf value
            tf_idf = dictionary[word]*(sentence.count(word)/len(sentence.split()))
            vector += (vec * tf_idf) # calculating tfidf weighted w2v
            tf_idf_weight += tf_idf
    if tf_idf_weight != 0:
        vector /= tf_idf_weight
    tfidf_w2v_vectors_test.append(vector)

print(len(tfidf_w2v_vectors_test))
print(len(tfidf_w2v_vectors_test[0]))
```

100%|██████████| 8250/8250 [00:27<00:00, 298.82it/s]

8250

300

In [73]:

```

tfidf_w2v_vectors_cv= []; # the avg-w2v for each sentence/review is stored in the
for sentence in tqdm(S_cv["clean_essays"]): # for each review/sentence
    vector = np.zeros(300) # as word vectors are of zero length
    tf_idf_weight = 0; # num of words with a valid vector in the sentence/review
    for word in sentence.split(): # for each word in a review/sentence
        if (word in glove_words) and (word in tfidf_words):
            vec = model[word] # getting the vector for each word
            # here we are multiplying idf value(dictionary[word]) and the tf value
            tf_idf = dictionary[word]*(sentence.count(word)/len(sentence.split()))
            vector += (vec * tf_idf) # calculating tfidf weighted w2v
            tf_idf_weight += tf_idf
    if tf_idf_weight != 0:
        vector /= tf_idf_weight
    tfidf_w2v_vectors_cv.append(vector)

print(len(tfidf_w2v_vectors_cv))
print(len(tfidf_w2v_vectors_cv[0]))

```

100%|██████████| 5025/5025 [00:16<00:00, 302.14it/s]

5025

300

In [74]:

```

# Similarly you can vectorize for title also
# average Word2Vec
# compute average word2vec for each review.
# S = ["abc def pqr", "def def def abc", "pqr pqr def"]
tfidf_model = TfidfVectorizer()
tfidf_model.fit(S_train["clean_pt"])
# we are converting a dictionary with word as a key, and the idf as a value
dictionary = dict(zip(tfidf_model.get_feature_names(), list(tfidf_model.idf_)))
tfidf_words = set(tfidf_model.get_feature_names())
tfidf_w2v_ppt_train= []; # the avg-w2v for each sentence/review is stored in this
for sentence in tqdm(S_train["clean_pt"]): # for each review/sentence
    vector = np.zeros(300) # as word vectors are of zero length
    tf_idf_weight = 0; # num of words with a valid vector in the sentence/review
    for word in sentence.split(): # for each word in a review/sentence
        if (word in glove_words) and (word in tfidf_words):
            vec = model[word] # getting the vector for each word
            # here we are multiplying idf value(dictionary[word]) and the tf value
            tf_idf = dictionary[word]*(sentence.count(word)/len(sentence.split()))
            vector += (vec * tf_idf) # calculating tfidf weighted w2v
            tf_idf_weight += tf_idf
    if tf_idf_weight != 0:
        vector /= tf_idf_weight
    tfidf_w2v_ppt_train.append(vector)

print(len(tfidf_w2v_ppt_train))
print(len(tfidf_w2v_ppt_train[0]))

```

100%|██████████| 11725/11725 [00:00<00:00, 19071.84it/s]

11725

300

```
In [75]: # Similarly you can vectorize for title also
# average Word2Vec
# compute average word2vec for each review.
# S = ["abc def pqr", "def def def abc", "pqr pqr def"]
tfidf_w2v_ppt_test= []; # the avg-w2v for each sentence/review is stored in this
for sentence in tqdm(S_test["clean_pt"]): # for each review/sentence
    vector = np.zeros(300) # as word vectors are of zero length
    tf_idf_weight = 0; # num of words with a valid vector in the sentence/review
    for word in sentence.split(): # for each word in a review/sentence
        if (word in glove_words) and (word in tfidf_words):
            vec = model[word] # getting the vector for each word
            # here we are multiplying idf value(dictionary[word]) and the tf value
            tf_idf = dictionary[word]*(sentence.count(word)/len(sentence.split()))
            vector += (vec * tf_idf) # calculating tfidf weighted w2v
            tf_idf_weight += tf_idf
    if tf_idf_weight != 0:
        vector /= tf_idf_weight
    tfidf_w2v_ppt_test.append(vector)

print(len(tfidf_w2v_ppt_test))
print(len(tfidf_w2v_ppt_test[0]))
```

100%|██████████| 8250/8250 [00:00<00:00, 18570.85it/s]

8250

300

```
In [76]: tfidf_w2v_ppt_cv= []; # the avg-w2v for each sentence/review is stored in this L
for sentence in tqdm(S_cv["clean_pt"]): # for each review/sentence
    vector = np.zeros(300) # as word vectors are of zero length
    tf_idf_weight = 0; # num of words with a valid vector in the sentence/review
    for word in sentence.split(): # for each word in a review/sentence
        if (word in glove_words) and (word in tfidf_words):
            vec = model[word] # getting the vector for each word
            # here we are multiplying idf value(dictionary[word]) and the tf value
            tf_idf = dictionary[word]*(sentence.count(word)/len(sentence.split()))
            vector += (vec * tf_idf) # calculating tfidf weighted w2v
            tf_idf_weight += tf_idf
    if tf_idf_weight != 0:
        vector /= tf_idf_weight
    tfidf_w2v_ppt_cv.append(vector)

print(len(tfidf_w2v_ppt_cv))
print(len(tfidf_w2v_ppt_cv[0]))
```

100%|██████████| 5025/5025 [00:00<00:00, 19367.77it/s]

5025

300

1.5.3 Vectorizing Numerical features

```
In [77]: price_data = resource_data.groupby('id').agg({'price': 'sum', 'quantity': 'sum'})
project_data = pd.merge(project_data, price_data, on='id', how='left')
```

```
In [78]: S_train = pd.merge(S_train, price_data, on='id', how='left')
S_test = pd.merge(S_test, price_data, on='id', how='left')
S_cv = pd.merge(S_cv, price_data, on='id', how='left')
```

Normalizing Price

```
In [79]: # check this one: https://www.youtube.com/watch?v=0H0q0cLn3Z4&t=530s
# standardization sklearn: https://scikit-learn.org/stable/modules/generated/sklearn.preprocessing import Normalizer

price_scalar = Normalizer()
X=S_train['price'].fillna(S_train['price'].mean())
Y=S_test['price'].fillna(S_test['price'].mean())
Z=S_cv['price'].fillna(S_cv['price'].mean())
price_scalar.fit(X.values.reshape(-1,1)) # finding the mean and standard deviation
price_standardized_train = price_scalar.transform(X.values.reshape(-1, 1))
price_standardized_test = price_scalar.transform(Y.values.reshape(-1, 1))
price_standardized_cv = price_scalar.transform(Z.values.reshape(-1, 1))
```

```
In [80]: print(price_standardized_train.shape)
print(price_standardized_test.shape)
print(price_standardized_cv.shape)
```

```
(11725, 1)
(8250, 1)
(5025, 1)
```

Normalizing number of previously posted projects

```
In [81]: price_scalar.fit(S_train['teacher_number_of_previously_posted_projects'].values)
prev_project_standardized_train = price_scalar.transform(S_train['teacher_number_of_previously_posted_projects'].values)
prev_project_standardized_test = price_scalar.transform(S_test['teacher_number_of_previously_posted_projects'].values)
prev_project_standardized_cv = price_scalar.transform(S_cv['teacher_number_of_previously_posted_projects'].values)
```

```
In [82]: print(prev_project_standardized_train.shape)
print(prev_project_standardized_test.shape)
print(prev_project_standardized_cv.shape)
```

```
(11725, 1)
(8250, 1)
(5025, 1)
```

Normalizing Quantity

In [83]:

```

X=S_train['quantity'].fillna(S_train['quantity'].mean())
Y=S_test['quantity'].fillna(S_test['quantity'].mean())
Z=S_cv['quantity'].fillna(S_cv['quantity'].mean())
price_scalar.fit(X.values.reshape(-1,1)) # finding the mean and standard deviation
quantity_standardized_train = price_scalar.transform(X.values.reshape(-1, 1))
quantity_standardized_test = price_scalar.transform(Y.values.reshape(-1, 1))
quantity_standardized_cv = price_scalar.transform(Z.values.reshape(-1, 1))

```

In [84]:

```

print(quantity_standardized_train.shape)
print(quantity_standardized_test.shape)
print(quantity_standardized_cv.shape)

```

```

(11725, 1)
(8250, 1)
(5025, 1)

```

normalizing title word count

In [85]:

```

normalizer = Normalizer()
normalizer.fit(S_train['title_word_count'].values.reshape(-1,1))
title_word_count_train = normalizer.transform(S_train['title_word_count'].values)
title_word_count_cv = normalizer.transform(S_cv['title_word_count'].values.reshape(-1,1))
title_word_count_test = normalizer.transform(S_test['title_word_count'].values.reshape(-1,1))
print("After vectorizations")
print(title_word_count_train.shape, y_train.shape)
print(title_word_count_cv.shape, y_cv.shape)
print(title_word_count_test.shape, y_test.shape)

```

```

After vectorizations
(11725, 1) (11725,)
(5025, 1) (5025,)
(8250, 1) (8250,)

```

NORMALIZING ESSAY WORD COUNT

In [86]:

```

normalizer = Normalizer()
normalizer.fit(S_train['essay_word_count'].values.reshape(-1,1))
essay_word_count_train = normalizer.transform(S_train['essay_word_count'].values)
essay_word_count_cv = normalizer.transform(S_cv['essay_word_count'].values.reshape(-1,1))
essay_word_count_test = normalizer.transform(S_test['essay_word_count'].values.reshape(-1,1))
print("After vectorizations")
print(essay_word_count_train.shape, y_train.shape)
print(essay_word_count_cv.shape, y_cv.shape)
print(essay_word_count_test.shape, y_test.shape)

```

```

After vectorizations
(11725, 1) (11725,)
(5025, 1) (5025,)
(8250, 1) (8250,)

```

```
In [87]: normalizer = Normalizer()
normalizer.fit(S_train['essay_word_count'].values.reshape(-1,1))
essay_word_count_train = normalizer.transform(S_train['essay_word_count'].values.reshape(-1,1))
essay_word_count_cv = normalizer.transform(S_cv['essay_word_count'].values.reshape(-1,1))
essay_word_count_test = normalizer.transform(S_test['essay_word_count'].values.reshape(-1,1))
print("After vectorizations")
print(essay_word_count_train.shape, y_train.shape)
print(essay_word_count_cv.shape, y_cv.shape)
print(essay_word_count_test.shape, y_test.shape)
```

```
After vectorizations
(11725, 1) (11725,)
(5025, 1) (5025,)
(8250, 1) (8250,)
```

NORMALIZING ESSAY SENTIMENT-POS

```
In [88]: normalizer = Normalizer()
normalizer.fit(S_train['pos'].values.reshape(-1,1))
essay_sent_pos_train = normalizer.transform(S_train['pos'].values.reshape(-1,1))
essay_sent_pos_cv = normalizer.transform(S_cv['pos'].values.reshape(-1,1))
essay_sent_pos_test = normalizer.transform(S_test['pos'].values.reshape(-1,1))
print("After vectorizations")
print(essay_sent_pos_train.shape, y_train.shape)
print(essay_sent_pos_cv.shape, y_cv.shape)
print(essay_sent_pos_test.shape, y_test.shape)
```

```
After vectorizations
(11725, 1) (11725,)
(5025, 1) (5025,)
(8250, 1) (8250,)
```

NORMALIZING ESSAY SENTIMENT-NEG

```
In [89]: normalizer = Normalizer()
normalizer.fit(S_train['neg'].values.reshape(-1,1))
essay_sent_neg_train = normalizer.transform(S_train['neg'].values.reshape(-1,1))
essay_sent_neg_cv = normalizer.transform(S_cv['neg'].values.reshape(-1,1))
essay_sent_neg_test = normalizer.transform(S_test['neg'].values.reshape(-1,1))
print("After vectorizations")
print(essay_sent_neg_train.shape, y_train.shape)
print(essay_sent_neg_cv.shape, y_cv.shape)
print(essay_sent_neg_test.shape, y_test.shape)
```

```
After vectorizations
(11725, 1) (11725,)
(5025, 1) (5025,)
(8250, 1) (8250,)
```

NORMALIZING ESSAY SENTIMENT-NEU

```
In [90]: normalizer = Normalizer()
normalizer.fit(S_train['neu'].values.reshape(-1,1))
essay_sent_neu_train = normalizer.transform(S_train['neu'].values.reshape(-1,1))
essay_sent_neu_cv = normalizer.transform(S_cv['neu'].values.reshape(-1,1))
essay_sent_neu_test = normalizer.transform(S_test['neu'].values.reshape(-1,1))
print("After vectorizations")
print(essay_sent_neu_train.shape, y_train.shape)
print(essay_sent_neu_cv.shape, y_cv.shape)
print(essay_sent_neu_test.shape, y_test.shape)
```

```
After vectorizations
(11725, 1) (11725,)
(5025, 1) (5025,)
(8250, 1) (8250,)
```

NORMALIZING ESSAY SENTIMEN-COMPOUND

```
In [91]: normalizer = Normalizer()
normalizer.fit(S_train['compound'].values.reshape(-1,1))
essay_sent_comp_train = normalizer.transform(S_train['compound'].values.reshape(-1,1))
essay_sent_comp_cv = normalizer.transform(S_cv['compound'].values.reshape(-1,1))
essay_sent_comp_test = normalizer.transform(S_test['compound'].values.reshape(-1,1))
print("After vectorizations")
print(essay_sent_comp_train.shape, y_train.shape)
print(essay_sent_comp_cv.shape, y_cv.shape)
print(essay_sent_comp_test.shape, y_test.shape)
```

```
After vectorizations
(11725, 1) (11725,)
(5025, 1) (5025,)
(8250, 1) (8250,)
```

1.5.4 Merging all the above features

- we need to merge all the numerical vectors i.e catogorical, text, numerical vectors

__ Computing Sentiment Scores __

Assignment 9: RF and GBDT

1. Apply both Random Forrest and GBDT on these feature sets

- **Set 1:** categorical (instead of one hot encoding, try [response coding](https://www.appliedaicourse.com/course/applied-ai-course-online/lessons/handling-categorical-and-numerical-features/) (<https://www.appliedaicourse.com/course/applied-ai-course-online/lessons/handling-categorical-and-numerical-features/>): use probability values), numerical features + project_title (BOW) + preprocessed_eassay (BOW)
- **Set 2:** categorical (instead of one hot encoding, try [response coding](https://www.appliedaicourse.com/course/applied-ai-course-online/lessons/handling-categorical-and-numerical-features/) (<https://www.appliedaicourse.com/course/applied-ai-course-online/lessons/handling-categorical-and-numerical-features/>))

[categorical-and-numerical-features/](#)): use probability values), numerical features + project_title(TFIDF)+ preprocessed_eassay (TFIDF)

- **Set 3:** categorical(instead of one hot encoding, try [response coding](#) (<https://www.appliedaiaicourse.com/course/applied-ai-course-online/lessons/handling-categorical-and-numerical-features/>): use probability values), numerical features + project_title(AVG W2V)+ preprocessed_eassay (AVG W2V)
- **Set 4:** categorical(instead of one hot encoding, try [response coding](#) (<https://www.appliedaiaicourse.com/course/applied-ai-course-online/lessons/handling-categorical-and-numerical-features/>): use probability values), numerical features + project_title(TFIDF W2V)+ preprocessed_eassay (TFIDF W2V)

2. The hyper paramter tuning (Consider any two hyper parameters preferably **n_estimators**, **max_depth**)

- Consider the following range for hyperparameters **n_estimators** = [10, 50, 100, 150, 200, 300, 500, 1000], **max_depth** = [2, 3, 4, 5, 6, 7, 8, 9, 10]
- Find the best hyper parameter which will give the maximum [AUC](#) (<https://www.appliedaiaicourse.com/course/applied-ai-course-online/lessons/receiver-operating-characteristic-curve-roc-curve-and-auc-1/>) value
- find the best hyper paramter using k-fold cross validation/simple cross validation data
- use gridsearch cv or randomsearch cv or you can write your own for loops to do this task

3. Representation of results

- You need to plot the performance of model both on train data and cross validation data for each hyper parameter, like shown in the figure



with X-axis as **n_estimators**, Y-axis as **max_depth**, and Z-axis as **AUC Score** , we have given the notebook which explains how to plot this 3d plot, you can find it in the same drive [3d_scatter_plot.ipynb](#)

or

- You need to plot the performance of model both on train data and cross validation data for each hyper parameter, like shown in the figure



[seaborn heat maps](#) (<https://seaborn.pydata.org/generated/seaborn.heatmap.html>) with rows as **n_estimators**, columns as **max_depth**, and values inside the cell representing **AUC Score**

- You can choose either of the plotting techniques: 3d plot or heat map
- Once after you found the best hyper parameter, you need to train your model with it, and find the AUC on test data and plot the ROC curve on both train and test.



- Along with plotting ROC curve, you need to print the [confusion matrix](#) (<https://www.appliedaiaicourse.com/course/applied-ai-course-online/lessons/confusion-matrix-tpr-fpr-fnr-tnr-1/>) with predicted and original labels of test data points



4. Conclusion

- You need to summarize the results at the end of the notebook, summarize it in the table format. To print out a table please refer to this prettytable library [link](http://zetcode.com/python/prettytable/) (<http://zetcode.com/python/prettytable/>)



Note: Data Leakage

1. There will be an issue of data-leakage if you vectorize the entire data and then split it into train/cv/test.
2. To avoid the issue of data-leakage, make sure to split your data first and then vectorize it.
3. While vectorizing your data, apply the method fit_transform() on you train data, and apply the method transform() on cv/test data.
4. For more details please go through this [link](https://soundcloud.com/applied-ai-course/leakage-bow-and-tfidf). (<https://soundcloud.com/applied-ai-course/leakage-bow-and-tfidf>)

Note:

1.I have used all the datapoints Random forest classifier for 3 vectorization techniques i.e BOW ,TFIDF and avgw2v 2.Since it was taking more than 2 days to compute i restarted my kernel and used 25K datapoints for Random Forest tfidf2v and for all 4 xgboost classifiers. 3.Response coding was again done on 25K points for train,test and cv separately.

Feature set 1 using BOW

```
In [92]: from scipy.sparse import hstack
```

```
In [93]: # merge two sparse matrices: https://stackoverflow.com/a/19710648/4084039

# with the same hstack function we are concatenating a sparse matrix and a dense
S_BOW_train= hstack((a1,b1,c1,d1,e1,text_bow,title_bow_train,price_standardized_
print(S_BOW_train.shape)

(11725, 29738)
```

```
In [94]: S_BOW_test= hstack((a2,b2,c2,d2,e2,text_bow_test,title_bow_test,price_standardized_
print(S_BOW_test.shape)

(8250, 29738)
```



```
In [95]: S_BOW_cv= hstack((a3,b3,c3,d3,e3,text_bow_cv,title_bow_cv,price_standardized_cv,
print(S_BOW_cv.shape)

(5025, 29738)
```

```
In [96]: def batch_predict(clf, data):
    # roc_auc_score(y_true, y_score) the 2nd parameter should be probability estimate
    # not the predicted outputs

    y_data_pred = []
    tr_loop = data.shape[0] - data.shape[0]%1000
    # consider you X_tr shape is 49041, then your tr_loop will be 49041 - 49041%1000
    # in this for loop we will iterate until the last 1000 multiplier
    for i in range(0, tr_loop, 1000):
        y_data_pred.extend(clf.predict_proba(data[i:i+1000])[:,1])
    # we will be predicting for the last data points
    if data.shape[0]%1000 !=0:
        y_data_pred.extend(clf.predict_proba(data[tr_loop:])[:,1])

    return y_data_pred
```

finding best hyperparameter using CV

```


In [153]: from sklearn.ensemble import RandomForestClassifier
import matplotlib.pyplot as plt
from sklearn.metrics import roc_auc_score
train_auc = []
cv_auc = []
a = []
b = []
import math
max_depth=[1, 5, 10, 50, 100, 500, 1000]

for i in tqdm(max_depth):
    rfc= RandomForestClassifier(max_depth=i,class_weight="balanced")
    l=rfc.fit(S_BOW_train, y_train)
    y_train_pred = batch_predict(rfc,S_BOW_train)
    y_cv_pred = batch_predict(rfc, S_BOW_cv)
    # roc_auc_score(y_true, y_score) the 2nd parameter should be probability estimate
    # not the predicted outputs
    train_auc.append(roc_auc_score(y_train,y_train_pred))
    cv_auc.append(roc_auc_score(y_cv, y_cv_pred))
    a.append(y_train_pred)
    b.append(y_cv_pred)

plt.plot([math.log(i) for i in max_depth],train_auc, label='Train AUC')
plt.plot([math.log(i) for i in max_depth],cv_auc, label='CV AUC')
plt.scatter([math.log(i) for i in max_depth],train_auc, label='Train AUC points')
plt.scatter([math.log(i) for i in max_depth],cv_auc, label='CV AUC points')
plt.legend()
plt.xlabel("max_depth-hyperparameter")
plt.ylabel("AUC")
plt.title("ERROR PLOTS")
plt.grid()
plt.show()

```

100%|██████████| 7/7 [01:47<00:00, 21.09s/it]



using Gridsearch CV for finding best hyperparameter

```


In [155]: from sklearn.ensemble import RandomForestClassifier
import matplotlib.pyplot as plt
from sklearn.metrics import roc_auc_score
train_auc = []
cv_auc = []
a = []
b = []
import math
n_estimators=[10,50,100,150,200,300,500,1000]

for i in tqdm(n_estimators):
    rfc= RandomForestClassifier(n_estimators=i,class_weight="balanced")
    l=rfc.fit(S_BOW_train, y_train)
    y_train_pred = batch_predict(rfc,S_BOW_train)
    y_cv_pred = batch_predict(rfc, S_BOW_cv)
    # roc_auc_score(y_true, y_score) the 2nd parameter should be probability estimates
    # not the predicted outputs
    train_auc.append(roc_auc_score(y_train,y_train_pred))
    cv_auc.append(roc_auc_score(y_cv, y_cv_pred))
    a.append(y_train_pred)
    b.append(y_cv_pred)

plt.plot([math.log(i) for i in n_estimators],train_auc, label='Train AUC')
plt.plot([math.log(i) for i in n_estimators],cv_auc, label='CV AUC')
plt.scatter([math.log(i) for i in n_estimators],train_auc, label='Train AUC points')
plt.scatter([math.log(i) for i in n_estimators],cv_auc, label='CV AUC points')
plt.legend()
plt.xlabel("n_estimators-hyperparameter")
plt.ylabel("AUC")
plt.title("ERROR PLOTS")
plt.grid()
plt.show()

```

100%|██████████| 8/8 [1:57:55<00:00, 1570.13s/it]



```

In [111]: from sklearn.ensemble import RandomForestClassifier

```

In [193]:

```

from sklearn.model_selection import GridSearchCV
from sklearn.metrics import confusion_matrix, auc, roc_auc_score, roc_curve

train_auc = []
cv_auc = []
train_auc_std = []
cv_auc_std = []

estimators = [10,50,100,150,200,300,500,1000]
depths = [2,3,4,5,6,7,8,9,10]

param_grid = {'n_estimators': estimators, 'max_depth':depths }
RFC = RandomForestClassifier()
model = GridSearchCV(RFC, param_grid, scoring = 'roc_auc', cv=3 , n_jobs = -1,pro
model.fit(S_BOW_train, y_train)
train_auc = model.cv_results_['mean_train_score']
train_auc_std = model.cv_results_['std_train_score']
cv_auc = model.cv_results_['mean_test_score']
cv_auc_std= model.cv_results_['std_test_score']

print("Model with best parameters :\n",model.best_estimator_)

```

Model with best parameters :

```

RandomForestClassifier(bootstrap=True, class_weight=None, criterion='gini',
                        max_depth=9, max_features='auto', max_leaf_nodes=None,
                        min_impurity_decrease=0.0, min_impurity_split=None,
                        min_samples_leaf=1, min_samples_split=2,
                        min_weight_fraction_leaf=0.0, n_estimators=1000, n_jobs=1,
                        oob_score=False, random_state=None, verbose=0,
                        warm_start=False)

```

In [194]:

```

best_depth = model.best_estimator_.max_depth
print(best_depth)

```

9

In [195]:

```


best_n_estimator = model.best_estimator_.n_estimators
print(best_n_estimator)

```


1000

```
In [196]: df_gridsearch = pd.DataFrame(model.cv_results_)
max_scores = df_gridsearch.groupby(['param_max_depth', 'param_n_estimators']).max
max_scores = max_scores.unstack()[['mean_test_score', 'mean_train_score']]
sns.heatmap(max_scores.mean_test_score, annot=True, fmt='.4g')
plt.title('AUC value on max_depth and e_estimators on CV data')
plt.show()
```

AUC value on max_depth and e_estimators on CV data



```
In [114]: #https://scikitlearn.org/stable/modules/generated/sklearn.metrics.roc_curve.html
from sklearn.metrics import roc_curve, auc
model = RandomForestClassifier(max_depth = 9,n_estimators=1000,random_state=0, c
model.fit(S_BOW_train, y_train)
# roc_auc_score(y_true, y_score) the 2nd parameter should be probability estimate
# not the predicted outputs
y_train_pred = batch_predict(model, S_BOW_train)
y_test_pred = batch_predict(model, S_BOW_test)
train_fpr, train_tpr, tr_thresholds = roc_curve(y_train, y_train_pred)
test_fpr, test_tpr, te_thresholds = roc_curve(y_test, y_test_pred)
plt.plot(train_fpr, train_tpr, label="Train AUC =" +str(auc(train_fpr, train_tpr))
plt.plot(test_fpr, test_tpr, label="Test AUC =" +str(auc(test_fpr, test_tpr)))
plt.legend()
plt.xlabel("True Positive Rate(TPR)")
plt.ylabel("False Positive Rate(FPR)")
plt.title("AUC")
plt.grid()
plt.show()
```




```
In [97]: def prediction(proba, threshold, fpr, tpr):
    t = threshold[np.argmax(fpr*(1-tpr))]
    # (tpr*(1-fpr)) will be maximum if your fpr is very low and tpr is very high
    print("the maximum value of tpr*(1-fpr)", max(tpr*(1-fpr)), "for threshold",
    predictions = []
    for i in proba:
        if i>=t:
            predictions.append(1)
        else:
            predictions.append(0)
    return predictions
```

confusion matrix for train data

```
In [116]: conf_matr_df_train = pd.DataFrame(confusion_matrix(y_train, prediction(y_train_pred,
train_fpr, train_tpr)), range(2),range(2))
sns.set(font_scale=1.4)#for label size
sns.heatmap(conf_matr_df_train, annot=True,annot_kws={"size": 16}, fmt='g')
```

the maximum value of $tpr*(1-fpr)$ 0.5293709542900648 for threshold 0.499

```
Out[116]: <matplotlib.axes._subplots.AxesSubplot at 0x2c14d08b2e8>
```




Confusion matrix for test data

```
In [117]: conf_matr_df_test = pd.DataFrame(confusion_matrix(y_test, prediction(y_test_pred,
train_fpr, train_tpr)), range(2),range(2))
sns.set(font_scale=1.4)#for label size
sns.heatmap(conf_matr_df_test, annot=True,annot_kws={"size": 16}, fmt='g')
```

the maximum value of $tpr*(1-fpr)$ 0.5293709542900648 for threshold 0.499

```
Out[117]: <matplotlib.axes._subplots.AxesSubplot at 0x2c12d0c8fd0>
```



Feature set 2 USING TFIDF_Train

```
In [98]: # Please write all the code with proper documentation  
from scipy.sparse import hstack  
# with the same hstack function we are concatenating a sparse matrix and a dense  
S_TFIDF_train= hstack((a1,b1,c1,d1,e1,text_tfidf_train,title_tfidf_train,price_s  
S_TFIDF_train.shape
```

Out[98]: (11725, 29738)

```
In [99]: S_TFIDF_test= hstack((a2,b2,c2,d2,e2,text_tfidf_test,title_tfidf_test,price_stand  
S_TFIDF_test.shape
```

Out[99]: (8250, 29738)

```
In [100]: S_TFIDF_cv= hstack((a3,b3,c3,d3,e3 ,text_tfidf_cv,title_tfidf_cv,price_standardi  
S_TFIDF_cv.shape
```

Out[100]: (5025, 29738)

Finding best parameter using CV


```
In [162]: from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import GridSearchCV
from sklearn.metrics import confusion_matrix, auc, roc_auc_score, roc_curve

train_auc = []
cv_auc = []
train_auc_std = []
cv_auc_std = []

estimators = [10,50,100,150,200,300,500,1000]
depths = [2,3,4,5,6,7,8,9,10]

param_grid = {'n_estimators': estimators, 'max_depth': depths }
RFC = RandomForestClassifier()
model1 = GridSearchCV(RFC, param_grid, scoring = 'roc_auc', cv=3 , n_jobs = -1, p
model1.fit(S_TFIDF_train, y_train)
train_auc = model1.cv_results_['mean_train_score']
train_auc_std = model1.cv_results_['std_train_score']
cv_auc = model1.cv_results_['mean_test_score']
cv_auc_std = model1.cv_results_['std_test_score']
print("Model with best parameters :\n", model1.best_estimator_)
```

Model with best parameters :

```
RandomForestClassifier(bootstrap=True, class_weight=None, criterion='gini',
                        max_depth=10, max_features='auto', max_leaf_nodes=None,
                        min_impurity_decrease=0.0, min_impurity_split=None,
                        min_samples_leaf=1, min_samples_split=2,
                        min_weight_fraction_leaf=0.0, n_estimators=1000, n_jobs=1,
                        oob_score=False, random_state=None, verbose=0,
                        warm_start=False)
```

Finding best hyperparameter using GridSearchCV

```
In [184]: best_depth = model1.best_estimator_.max_depth
print(best_depth)
```


10

```
In [186]: best_estimator = model1.best_estimator_.n_estimators
print(best_estimator)
```

1000


```
In [190]: df_gridsearch = pd.DataFrame(model1.cv_results_)
max_scores = df_gridsearch.groupby(['param_max_depth', 'param_n_estimators']).max
max_scores = max_scores.unstack()[['mean_test_score', 'mean_train_score']]
sns.heatmap(max_scores.mean_test_score, annot=True, fmt='.4g')
plt.title('AUC value on max_depth and e_estimators on CV data')
plt.show()
```

AUC value on max_depth and e_estimators on CV data



Max_AUC=0.66,Max_depth=10,n_estimators=1000

```
In [187]: #https://scikitlearn.org/stable/modules/generated/sklearn.metrics.roc_curve.html
from sklearn.metrics import roc_curve, auc
model = RandomForestClassifier(max_depth=10,n_estimators=1000,random_state=0, class_weight='balanced')
model.fit(S_TFIDF_train, y_train)
# roc_auc_score(y_true, y_score) the 2nd parameter should be probability estimates
# not the predicted outputs
y_train_pred = batch_predict(model1, S_TFIDF_train)
y_test_pred = batch_predict(model1, S_TFIDF_test)
train_fpr, train_tpr, tr_thresholds = roc_curve(y_train, y_train_pred)
test_fpr, test_tpr, te_thresholds = roc_curve(y_test, y_test_pred)
plt.plot(train_fpr, train_tpr, label="Train AUC =" + str(auc(train_fpr, train_tpr)))
plt.plot(test_fpr, test_tpr, label="Test AUC =" + str(auc(test_fpr, test_tpr)))
plt.legend()
plt.xlabel("True Positive Rate(TPR)")
plt.ylabel("False Positive Rate(FPR)")
plt.title("AUC")
plt.grid()
plt.show()
```




confusion matrix for train data

```
In [188]: conf_matr_df_train = pd.DataFrame(confusion_matrix(y_train, prediction(y_train_pred,
train_fpr, train_tpr)), range(2),range(2))
sns.set(font_scale=1.4)#for label size
sns.heatmap(conf_matr_df_train, annot=True,annot_kws={"size": 16}, fmt='g')
```

the maximum value of $tpr*(1-fpr)$ 0.5805688843892902 for threshold 0.504

Out[188]: <matplotlib.axes._subplots.AxesSubplot at 0x1b7e24b9390>




Confusion matrix for test data

```
In [189]: conf_matr_df_test = pd.DataFrame(confusion_matrix(y_test, prediction(y_test_pred,
train_fpr, train_tpr)), range(2),range(2))
sns.set(font_scale=1.4)#for label size
sns.heatmap(conf_matr_df_test, annot=True,annot_kws={"size": 16}, fmt='g')
```

the maximum value of $tpr*(1-fpr)$ 0.5805688843892902 for threshold 0.504

Out[189]: <matplotlib.axes._subplots.AxesSubplot at 0x1b7804c1048>



Feature set 3 USING AVG_W2V

```
In [101]: from scipy.sparse import hstack  
# with the same hstack function we are concatenating a sparse matrix and a dense  
S_avgw2v_train= hstack((a1,b1,c1,d1,e1,avg_w2v_vectors_train,avg_w2v_title_train,  
print(S_avgw2v_train.shape)
```

(11725, 619)

```
In [102]: S_avgw2v_test= hstack((a2,b2,c2,d2,e2,avg_w2v_vectors_test,avg_w2v_title_test,pr  
print(S_avgw2v_test.shape)
```

(8250, 619)

```
In [103]: S_avgw2v_cv= hstack((a3,b3,c3,d3,e3,avg_w2v_vectors_cv,avg_w2v_title_cv,price_st  
print(S_avgw2v_cv.shape)
```

(5025, 619)

FINDING BEST HYPERPARAMETER USING CV

FINDING BEST HYPERPARAMETER USING GRIDSEARCHCV

```
In [121]: from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import GridSearchCV
from sklearn.metrics import confusion_matrix, auc, roc_auc_score, roc_curve

train_auc = []
cv_auc = []
train_auc_std = []
cv_auc_std = []

estimators = [10,50,100,150,200,300,500,1000]
depths = [2,3,4,5,6,7,8,9,10]

param_grid = {'n_estimators': estimators, 'max_depth': depths }
RFC = RandomForestClassifier()
model2 = GridSearchCV(RFC, param_grid, scoring = 'roc_auc', cv=3 , n_jobs = -1,p
model2.fit(S_avgw2v_train, y_train)
train_auc = model2.cv_results_['mean_train_score']
train_auc_std = model2.cv_results_['std_train_score']
cv_auc = model2.cv_results_['mean_test_score']
cv_auc_std= model2.cv_results_['std_test_score']


print("Model with best parameters :\n",model2.best_estimator_)
```

Model with best parameters :


```
RandomForestClassifier(bootstrap=True, class_weight=None, criterion='gini',
                        max_depth=10, max_features='auto', max_leaf_nodes=None,
                        min_impurity_decrease=0.0, min_impurity_split=None,
                        min_samples_leaf=1, min_samples_split=2,
                        min_weight_fraction_leaf=0.0, n_estimators=1000, n_jobs=1,
                        oob_score=False, random_state=None, verbose=0,
                        warm_start=False)
```

```
In [129]: df_gridsearch = pd.DataFrame(model2.cv_results_)
max_scores = df_gridsearch.groupby(['param_max_depth', 'param_n_estimators']).max
max_scores = max_scores.unstack()[['mean_test_score', 'mean_train_score']]
sns.heatmap(max_scores.mean_test_score, annot=True, fmt='.4g')
plt.title('AUC value on max_depth and e_estimators on CV data')
plt.show()
```

AUC value on max_depth and e_estimators on CV data



```
In [126]: #https://scikitlearn.org/stable/modules/generated/sklearn.metrics.roc_curve.html
from sklearn.metrics import roc_curve, auc
model3 = RandomForestClassifier(max_depth = 10,n_estimators=1000,random_state=0,
model3.fit(S_avgw2v_train, y_train)
# roc_auc_score(y_true, y_score) the 2nd parameter should be probability estimate
# not the predicted outputs
y_train_pred = batch_predict(model3, S_avgw2v_train)
y_test_pred = batch_predict(model3, S_avgw2v_test)
train_fpr, train_tpr, tr_thresholds = roc_curve(y_train, y_train_pred)
test_fpr, test_tpr, te_thresholds = roc_curve(y_test, y_test_pred)
plt.plot(train_fpr, train_tpr, label="Train AUC =" +str(auc(train_fpr, train_tpr))
plt.plot(test_fpr, test_tpr, label="Test AUC =" +str(auc(test_fpr, test_tpr)))
plt.legend()
plt.xlabel("True Positive Rate(TPR)")
plt.ylabel("False Positive Rate(FPR)")
plt.title("AUC")
plt.grid()
plt.show()
```




CONFUSION MATRIX FOR TRAIN DATA

```
In [127]: conf_matr_df_train = pd.DataFrame(confusion_matrix(y_train, prediction(y_train_proba_train_fpr, train_tpr)), range(2),range(2))
sns.set(font_scale=1.4)#for label size
sns.heatmap(conf_matr_df_train, annot=True,annot_kws={"size": 16}, fmt='g')
```

the maximum value of $tpr \cdot (1 - fpr)$ 0.8155078270088588 for threshold 0.52

```
Out[127]: <matplotlib.axes._subplots.AxesSubplot at 0x2c14cffffa90>
```




CONFUSION MATRIX FOR TEST DATA


```
In [128]: conf_matr_df_test = pd.DataFrame(confusion_matrix(y_test, prediction(y_test_pred,
train_fpr, train_tpr)), range(2),range(2))
sns.set(font_scale=1.4)#for label size
sns.heatmap(conf_matr_df_test, annot=True,annot_kws={"size": 16}, fmt='g')
```

the maximum value of $tpr \cdot (1 - fpr)$ 0.8155078270088588 for threshold 0.52

```
Out[128]: <matplotlib.axes._subplots.AxesSubplot at 0x2c14cebdc18>
```



FEATURE SET 4:TFIDF_W2V

```
In [104]: from scipy.sparse import hstack
# with the same hstack function we are concatenating a sparse matrix and a dense
S_tfidf_w2v_train= hstack((a1,b1,c1,d1,e1,tfidf_w2v_vectors_train,tfidf_w2v_ppt_train))
print(S_tfidf_w2v_train.shape)

S_tfidf_w2v_test= hstack((a2,b2,c2,d2,e2,tfidf_w2v_vectors_test,tfidf_w2v_ppt_test))
print(S_tfidf_w2v_test.shape)

S_tfidf_w2v_cv= hstack((a3,b3,c3,d3,e3,tfidf_w2v_vectors_cv,tfidf_w2v_ppt_cv,prior_cv))
print(S_tfidf_w2v_cv.shape)
```

```
(11725, 619)
(8250, 619)
(5025, 619)
```

Using CV to find best hyperparameter

```
In [120]: from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import GridSearchCV
from sklearn.metrics import confusion_matrix, auc, roc_auc_score, roc_curve

train_auc = []
cv_auc = []
train_auc_std = []
cv_auc_std = []

estimators = [10,50,100,150,200,300,500,1000]
depths = [2,3,4,5,6,7,8,9,10]

param_grid = {'n_estimators': estimators, 'max_depth': depths }
RFC = RandomForestClassifier()
model3 = GridSearchCV(RFC, param_grid, scoring = 'roc_auc', cv=3 , n_jobs = -1, parallel=True)
model3.fit(S_tfidf_w2v_train, y_train)
train_auc = model3.cv_results_['mean_train_score']
train_auc_std = model3.cv_results_['std_train_score']
cv_auc = model3.cv_results_['mean_test_score']
cv_auc_std = model3.cv_results_['std_test_score']

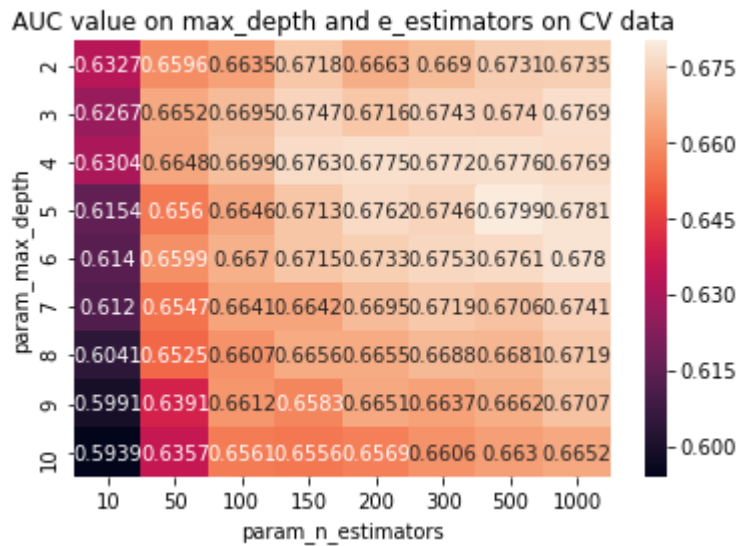
print("Model with best parameters :\n", model3.best_estimator_)
```

Model with best parameters :


```
RandomForestClassifier(bootstrap=True, class_weight=None, criterion='gini',
                        max_depth=5, max_features='auto', max_leaf_nodes=None,
                        min_impurity_decrease=0.0, min_impurity_split=None,
                        min_samples_leaf=1, min_samples_split=2,
                        min_weight_fraction_leaf=0.0, n_estimators=500, n_jobs=1,
                        oob_score=False, random_state=None, verbose=0,
                        warm_start=False)
```

Using GridsearchCV to find best hyperparameter

```
In [121]: df_gridsearch = pd.DataFrame(model3.cv_results_)
max_scores = df_gridsearch.groupby(['param_max_depth', 'param_n_estimators']).max
max_scores = max_scores.unstack()[['mean_test_score', 'mean_train_score']]
sns.heatmap(max_scores.mean_test_score, annot=True, fmt='.4g')
plt.title('AUC value on max_depth and e_estimators on CV data')
plt.show()
```



```
In [123]: #https://scikitlearn.org/stable/modules/generated/sklearn.metrics.roc_curve.html
from sklearn.metrics import roc_curve, auc
model4 = RandomForestClassifier(max_depth = 5,n_estimators=500,random_state=0, c
model4.fit(S_tfidf_w2v_train, y_train)
# roc_auc_score(y_true, y_score) the 2nd parameter should be probability estimate
# not the predicted outputs
y_train_pred = batch_predict(model4,S_tfidf_w2v_train)
y_test_pred = batch_predict(model4, S_tfidf_w2v_test)
train_fpr, train_tpr, tr_thresholds = roc_curve(y_train, y_train_pred)
test_fpr, test_tpr, te_thresholds = roc_curve(y_test, y_test_pred)
plt.plot(train_fpr, train_tpr, label="Train AUC =" +str(auc(train_fpr, train_tpr))
plt.plot(test_fpr, test_tpr, label="Test AUC =" +str(auc(test_fpr, test_tpr)))
plt.legend()
plt.xlabel("True Positive Rate(TPR)")
plt.ylabel("False Positive Rate(FPR)")
plt.title("AUC")
plt.grid()
plt.show()
```




confusion matrix for train data

```
In [124]: conf_matr_df_train = pd.DataFrame(confusion_matrix(y_train, prediction(y_train_proba_train_fpr, train_tpr)), range(2),range(2))
sns.set(font_scale=1.4)#for label
sns.heatmap(conf_matr_df_train, annot=True,annot_kws={"size": 16}, fmt='g')
```

the maximum value of $tpr \cdot (1 - fpr)$ 0.5699321417071109 for threshold 0.491

Out[124]: <matplotlib.axes._subplots.AxesSubplot at 0x1318b0baa58>




CONFUSION MATRIX FOR TEST DATA

```
In [125]: conf_matr_df_test = pd.DataFrame(confusion_matrix(y_test, prediction(y_test_pred,
train_fpr, train_tpr)), range(2),range(2))
sns.set(font_scale=1.4)#for label size
sns.heatmap(conf_matr_df_test, annot=True,annot_kws={"size": 16}, fmt='g')
```

the maximum value of $tpr \cdot (1 - fpr)$ 0.5699321417071109 for threshold 0.491

```
Out[125]: <matplotlib.axes._subplots.AxesSubplot at 0x1318b175eb8>
```



XGBOOST with BOW

```
In [107]: import xgboost as xgb
from sklearn.model_selection import GridSearchCV
from sklearn.metrics import confusion_matrix, auc, roc_auc_score, roc_curve
from xgboost import XGBClassifier
```

In [127]:

```

estimators = [10,50,100,150,200,300,500,1000]
Depths = [2,3,4,5,6,7,8,9,10]

param_grid = {'n_estimators': estimators, 'max_depth':Depths }
XGB = XGBClassifier(booster='gbtree')
xgb1 = GridSearchCV(XGB, param_grid, scoring = 'roc_auc', cv=3 , n_jobs = -1,pre
xgb1.fit(S_BOW_train, y_train)

```

```

Out[127]: GridSearchCV(cv=3, error_score='raise',
    estimator=XGBClassifier(base_score=0.5, booster='gbtree', colsample_bye
vel=1,
    colsample_bynode=1, colsample_bytree=1, gamma=0, learning_rate=0.1,
    max_delta_step=0, max_depth=3, min_child_weight=1, missing=None,
    n_estimators=100, n_jobs=1, nthread=None,
    objective='binary:logistic', random_state=0, reg_alpha=0,
    reg_lambda=1, scale_pos_weight=1, seed=None, silent=None,
    subsample=1, verbosity=1),
    fit_params=None, iid=True, n_jobs=-1,
    param_grid={'n_estimators': [10, 50, 100, 150, 200, 300, 500, 1000], 'ma
x_depth': [2, 3, 4, 5, 6, 7, 8, 9, 10]},
    pre_dispatch=2, refit=True, return_train_score='warn',
    scoring='roc_auc', verbose=0)

```


In [132]:

```

df_gridsearch1 = pd.DataFrame(xgb1.cv_results_)
max_scores1 = df_gridsearch1.groupby(['param_max_depth', 'param_n_estimators']).m
max_scores1 = max_scores1.unstack()[['mean_test_score', 'mean_train_score']]
sns.heatmap(max_scores1.mean_train_score, annot=True, fmt='.4g')
plt.title('AUC value on max_depth and e_estimators on Train data')
plt.show()


```

AUC value on max_depth and e_estimators on Train data



from heat map we can see we get max AUC value at best depth=3 and n_estimator=100


```
In [108]: from sklearn.metrics import roc_curve, auc
clf_1 = XGBClassifier(booster='gbtree',max_depth=3,n_estimators=100)
clf_1.fit(S_BOW_train,y_train)
# roc_auc_score(y_true, y_score) the 2nd parameter should be probability estimate
# not the predicted outputs
y_train_pred = batch_predict(clf_1,S_BOW_train)
y_test_pred = batch_predict(clf_1, S_BOW_test)
train_fpr, train_tpr, tr_thresholds = roc_curve(y_train, y_train_pred)
test_fpr, test_tpr, te_thresholds = roc_curve(y_test, y_test_pred)
plt.plot(train_fpr, train_tpr, label="Train AUC =" +str(auc(train_fpr, train_tpr)))
plt.plot(test_fpr, test_tpr, label="Test AUC =" +str(auc(test_fpr, test_tpr)))
plt.legend()
plt.xlabel("True Positive Rate(TPR)")
plt.ylabel("False Positive Rate(FPR)")
plt.title("AUC")
plt.grid()
plt.show()
```




```
In [109]: conf_matr_df_train = pd.DataFrame(confusion_matrix(y_train, prediction(y_train_pred,
train_fpr, train_tpr)), range(2),range(2))
sns.set(font_scale=1.4)#for label
sns.heatmap(conf_matr_df_train, annot=True,annot_kws={"size": 16}, fmt='g')
```

the maximum value of $tpr*(1-fpr)$ 0.5452355467382929 for threshold 0.835


```
Out[109]: <matplotlib.axes._subplots.AxesSubplot at 0x20cc0d9c080>
```



```
In [110]: conf_matr_df_test = pd.DataFrame(confusion_matrix(y_test, prediction(y_test_pred,
train_fpr, train_tpr)), range(2),range(2))
sns.set(font_scale=1.4)#for label size
sns.heatmap(conf_matr_df_test, annot=True,annot_kws={"size": 16}, fmt='g')
```

the maximum value of $tpr*(1-fpr)$ 0.5452355467382929 for threshold 0.835

```
Out[110]: <matplotlib.axes._subplots.AxesSubplot at 0x20cc0bcf2e8>
```

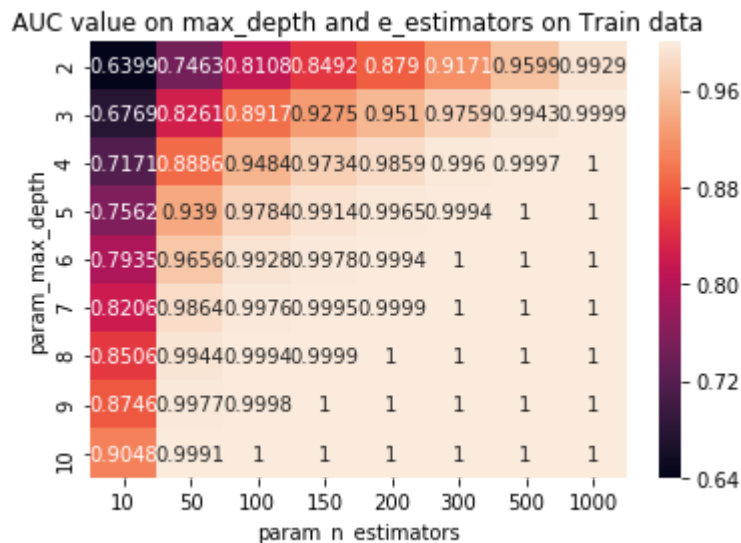


XGBOOST FOR TFIDF

```
In [108]: estimators = [10,50,100,150,200,300,500,1000]
Depths = [2,3,4,5,6,7,8,9,10]
param_grid = {'n_estimators': estimators, 'max_depth':Depths }
XGB = XGBClassifier(booster='gbtree')
xgb2 = GridSearchCV(XGB, param_grid, scoring = 'roc_auc', cv=3 , n_jobs = -1,pre
xgb2.fit(S_TFIDF_train, y_train)
```


```
Out[108]: GridSearchCV(cv=3, error_score='raise',
    estimator=XGBClassifier(base_score=0.5, booster='gbtree', colsample_byle
vel=1,
    colsample_bynode=1, colsample_bytree=1, gamma=0, learning_rate=0.1,
    max_delta_step=0, max_depth=3, min_child_weight=1, missing=None,
    n_estimators=100, n_jobs=1, nthread=None,
    objective='binary:logistic', random_state=0, reg_alpha=0,
    reg_lambda=1, scale_pos_weight=1, seed=None, silent=None,
    subsample=1, verbosity=1),
    fit_params=None, iid=True, n_jobs=-1,
    param_grid={'n_estimators': [10, 50, 100, 150, 200, 300, 500, 1000], 'ma
x_depth': [2, 3, 4, 5, 6, 7, 8, 9, 10]},
    pre_dispatch=2, refit=True, return_train_score='warn',
    scoring='roc_auc', verbose=0)
```

```
In [111]: df_gridsearch1 = pd.DataFrame(xgb2.cv_results_)
max_scores1 = df_gridsearch1.groupby(['param_max_depth', 'param_n_estimators']).ma
max_scores1 = max_scores1.unstack()[['mean_test_score', 'mean_train_score']]
sns.heatmap(max_scores1.mean_train_score, annot=True, fmt='.4g')
plt.title('AUC value on max_depth and e_estimators on Train data')
plt.show()
```



from heat map we can see we get max AUC value at best depth=3 and n_estimator=100

```
In [111]: #https://scikitlearn.org/stable/modules/generated/sklearn.metrics.roc_curve.html
from sklearn.metrics import roc_curve, auc
clf_2 = XGBClassifier(booster='gbtree',max_depth=3,n_estimators=100)
clf_2.fit(S_TFIDF_train,y_train)
# roc_auc_score(y_true, y_score) the 2nd parameter should be probability estimate
# not the predicted outputs
y_train_pred = batch_predict(clf_2,S_TFIDF_train)
y_test_pred = batch_predict(clf_2, S_TFIDF_test)
train_fpr, train_tpr, tr_thresholds = roc_curve(y_train, y_train_pred)
test_fpr, test_tpr, te_thresholds = roc_curve(y_test, y_test_pred)
plt.plot(train_fpr, train_tpr, label="Train AUC =" +str(auc(train_fpr, train_tpr)))
plt.plot(test_fpr, test_tpr, label="Test AUC =" +str(auc(test_fpr, test_tpr)))
plt.legend()
plt.xlabel("True Positive Rate(TPR)")
plt.ylabel("False Positive Rate(FPR)")
plt.title("AUC")
plt.grid()
plt.show()
```




Confusion matrix for train data

```
In [112]: conf_matr_df_train = pd.DataFrame(confusion_matrix(y_train, prediction(y_train_proba_train_fpr, train_tpr)), range(2),range(2))
sns.set(font_scale=1.4)#for label
sns.heatmap(conf_matr_df_train, annot=True,annot_kws={"size": 16}, fmt='g')
```

the maximum value of $tpr \cdot (1 - fpr)$ 0.5913520998720945 for threshold 0.831

```
Out[112]: <matplotlib.axes._subplots.AxesSubplot at 0x20cc0fec080>
```




Confusion matrix for test data

```
In [113]: conf_matr_df_test = pd.DataFrame(confusion_matrix(y_test, prediction(y_test_pred,
train_fpr, train_tpr)), range(2),range(2))
sns.set(font_scale=1.4)#for label size
sns.heatmap(conf_matr_df_test, annot=True,annot_kws={"size": 16}, fmt='g')
```

the maximum value of $tpr \cdot (1 - fpr)$ 0.5913520998720945 for threshold 0.831

Out[113]: <matplotlib.axes._subplots.AxesSubplot at 0x20cc1905550>



XGBOOST FOR AVG_W2V

```
In [ ]: Grid search
```

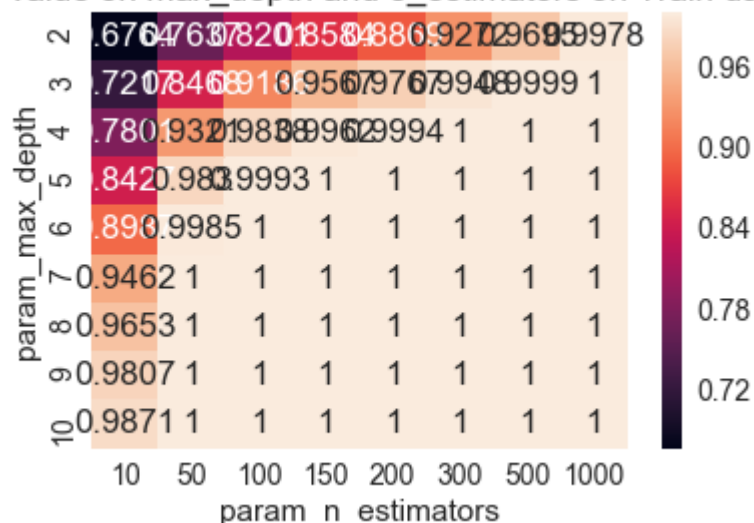
```
In [112]: estimators = [10,50,100,150,200,300,500,1000]
Depths = [2,3,4,5,6,7,8,9,10]
param_grid = {'n_estimators': estimators, 'max_depth':Depths }
XGB = XGBClassifier(booster='gbtree')
xgb3 = GridSearchCV(XGB, param_grid, scoring = 'roc_auc', cv=3 , n_jobs = -1,pre
xgb3.fit(S_avgw2v_train, y_train)
```

```
Out[112]: GridSearchCV(cv=3, error_score='raise',
      estimator=XGBClassifier(base_score=0.5, booster='gbtree', colsample_bye
vel=1,
      colsample_bynode=1, colsample_bytree=1, gamma=0, learning_rate=0.1,
      max_delta_step=0, max_depth=3, min_child_weight=1, missing=None,
      n_estimators=100, n_jobs=1, nthread=None,
      objective='binary:logistic', random_state=0, reg_alpha=0,
      reg_lambda=1, scale_pos_weight=1, seed=None, silent=None,
      subsample=1, verbosity=1),
      fit_params=None, iid=True, n_jobs=-1,
      param_grid={'n_estimators': [10, 50, 100, 150, 200, 300, 500, 1000], 'ma
x_depth': [2, 3, 4, 5, 6, 7, 8, 9, 10]},
      pre_dispatch=2, refit=True, return_train_score='warn',
      scoring='roc_auc', verbose=0)
```

Heat map


```
In [122]: df_gridsearch1 = pd.DataFrame(xgb3.cv_results_)
max_scores1 = df_gridsearch1.groupby(['param_max_depth','param_n_estimators']).m
max_scores1 = max_scores1.unstack()['mean_test_score', 'mean_train_score']]
sns.heatmap(max_scores1.mean_train_score, annot=True, fmt='.4g')
plt.title('AUC value on max_depth and e_estimators on Train data')
plt.show()
```

AUC value on max_depth and e_estimators on Train data



from heat map we can see we get max AUC value at best depth=3 and n_estimator=100

```
In [114]: #https://scikitlearn.org/stable/modules/generated/sklearn.metrics.roc_curve.html
from sklearn.metrics import roc_curve, auc
clf_4 = XGBClassifier(booster='gbtree',max_depth=3,n_estimators=100)
clf_4.fit(S_avgw2v_train,y_train)
# roc_auc_score(y_true, y_score) the 2nd parameter should be probability estimate
# not the predicted outputs
y_train_pred = batch_predict(clf_4,S_avgw2v_train)
y_test_pred = batch_predict(clf_4,S_avgw2v_test)
train_fpr, train_tpr, tr_thresholds = roc_curve(y_train, y_train_pred)
test_fpr, test_tpr, te_thresholds = roc_curve(y_test, y_test_pred)
plt.plot(train_fpr, train_tpr, label="Train AUC =" +str(auc(train_fpr, train_tpr)))
plt.plot(test_fpr, test_tpr, label="Test AUC =" +str(auc(test_fpr, test_tpr)))
plt.legend()
plt.xlabel("True Positive Rate(TPR)")
plt.ylabel("False Positive Rate(FPR)")
plt.title("AUC")
plt.grid()
plt.show()
```




Confusion matrix for train data

```
In [115]: conf_matr_df_train = pd.DataFrame(confusion_matrix(y_train, prediction(y_train_proba_train_fpr, train_tpr)), range(2),range(2))
sns.set(font_scale=1.4)#for label
sns.heatmap(conf_matr_df_train, annot=True,annot_kws={"size": 16}, fmt='g')
```

the maximum value of $tpr*(1-fpr)$ 0.6156967924024537 for threshold 0.848

Out[115]: <matplotlib.axes._subplots.AxesSubplot at 0x20cc1927828>




Confusion matrix for test data


```
In [116]: conf_matr_df_test = pd.DataFrame(confusion_matrix(y_test, prediction(y_test_pred,
train_fpr, train_tpr)), range(2),range(2))
sns.set(font_scale=1.4)#for label size
sns.heatmap(conf_matr_df_test, annot=True,annot_kws={"size": 16}, fmt='g')
```

the maximum value of $tpr \cdot (1 - fpr)$ 0.6156967924024537 for threshold 0.848

Out[116]: <matplotlib.axes._subplots.AxesSubplot at 0x20cc1979ba8>



XGBOOST FOR TFIDF_W2V


```
In [133]: estimators = [10,50,100,150,200,300,500,1000]
Depths = [2,3,4,5,6,7,8,9,10]
param_grid = {'n_estimators': estimators, 'max_depth':Depths }
XGB = XGBClassifier(booster='gbtree')
xgb4 = GridSearchCV(XGB, param_grid, scoring = 'roc_auc', cv=3 , n_jobs = -1,pre
xgb4.fit(S_tfidf_w2v_train, y_train)
```

```
Out[133]: GridSearchCV(cv=3, error_score='raise',
      estimator=XGBClassifier(base_score=0.5, booster='gbtree', colsample_bye
vel=1,
      colsample_bynode=1, colsample_bytree=1, gamma=0, learning_rate=0.1,
      max_delta_step=0, max_depth=3, min_child_weight=1, missing=None,
      n_estimators=100, n_jobs=1, nthread=None,
      objective='binary:logistic', random_state=0, reg_alpha=0,
      reg_lambda=1, scale_pos_weight=1, seed=None, silent=None,
      subsample=1, verbosity=1),
      fit_params=None, iid=True, n_jobs=-1,
      param_grid={'n_estimators': [10, 50, 100, 150, 200, 300, 500, 1000], 'ma
x_depth': [2, 3, 4, 5, 6, 7, 8, 9, 10]},
      pre_dispatch=2, refit=True, return_train_score='warn',
      scoring='roc_auc', verbose=0)
```

PERFORMING GRID SEARCH


```
In [134]: df_gridsearch1 = pd.DataFrame(xgb4.cv_results_)
max_scores1 = df_gridsearch1.groupby(['param_max_depth','param_n_estimators']).m
max_scores1 = max_scores1.unstack()['mean_test_score', 'mean_train_score']]
sns.heatmap(max_scores1.mean_train_score, annot=True, fmt='.4g')
plt.title('AUC value on max_depth and e_estimators on Train data')
plt.show()
```

AUC value on max_depth and e_estimators on Train data



from heat map we can see we get max AUC value at best depth=3 and n_estimator=100

```
In [117]: #https://scikitlearn.org/stable/modules/generated/sklearn.metrics.roc_curve.html
from sklearn.metrics import roc_curve, auc
clf_4 = XGBClassifier(booster='gbtree',max_depth=3,n_estimators=100)
clf_4.fit(S_tfidf_w2v_train,y_train)
# roc_auc_score(y_true, y_score) the 2nd parameter should be probability estimate
# not the predicted outputs
y_train_pred = batch_predict(clf_4, S_tfidf_w2v_train)
y_test_pred = batch_predict(clf_4, S_tfidf_w2v_test)
train_fpr, train_tpr, tr_thresholds = roc_curve(y_train, y_train_pred)
test_fpr, test_tpr, te_thresholds = roc_curve(y_test, y_test_pred)
plt.plot(train_fpr, train_tpr, label="Train AUC =" +str(auc(train_fpr, train_tpr)))
plt.plot(test_fpr, test_tpr, label="Test AUC =" +str(auc(test_fpr, test_tpr)))
plt.legend()
plt.xlabel("True Positive Rate(TPR)")
plt.ylabel("False Positive Rate(FPR)")
plt.title("AUC")
plt.grid()
plt.show()
```




CONFUSION MATRIX FOR TRAIN DATA

```
In [118]: conf_matr_df_train = pd.DataFrame(confusion_matrix(y_train, prediction(y_train_proba_train_fpr, train_tpr)), range(2),range(2))
sns.set(font_scale=1.4)#for label
sns.heatmap(conf_matr_df_train, annot=True,annot_kws={"size": 16}, fmt='g')
```

the maximum value of $tpr \cdot (1 - fpr)$ 0.6215320531859162 for threshold 0.837

Out[118]: <matplotlib.axes._subplots.AxesSubplot at 0x20cc1997898>




CONFUSION MATRIX FOR TEST DATA

```
In [119]: conf_matr_df_test = pd.DataFrame(confusion_matrix(y_test, prediction(y_test_pred,
train_fpr, train_tpr)), range(2),range(2))
sns.set(font_scale=1.4)#for label size
sns.heatmap(conf_matr_df_test, annot=True,annot_kws={"size": 16}, fmt='g')
```

the maximum value of $tpr \cdot (1 - fpr)$ 0.6215320531859162 for threshold 0.837

Out[119]: <matplotlib.axes._subplots.AxesSubplot at 0x20cc19d0fd0>



```
In [120]: from prettytable import PrettyTable
p= PrettyTable()
v1 = "BOW"
v2 = "TFIDF"
v3 = "AVG-W2V"
v4 = "TFIDF-W2V"
m1 = 'Random Forest'
m2 = 'GBDT-XGBOOST'
p.field_names = ["Vectorizer", "Model", "best_Depth", "base_learners", " AUC"]
p.add_row([v1,m1,9,1000,0.6677])
p.add_row([v2,m1,10,1000,0.6647])
p.add_row([v3,m1,10,1000,0.6672])
p.add_row([v4,m1,5,500,0.6799])
p.add_row([v1,m2,3,100,0.7858])
p.add_row([v2,m2,3,100,0.8917])
p.add_row([v3,m2,3,100,0.8913])
p.add_row([v4,m2,3,100,0.7913])
print(p)
```

Vectorizer	Model	best_Depth	base_learners	AUC
BOW	Random Forest	9	1000	0.6677
TFIDF	Random Forest	10	1000	0.6647
AVG-W2V	Random Forest	10	1000	0.6672
TFIDF-W2V	Random Forest	5	500	0.6799
BOW	GBDT-XGBOOST	3	100	0.7858
TFIDF	GBDT-XGBOOST	3	100	0.8917
AVG-W2V	GBDT-XGBOOST	3	100	0.8913
TFIDF-W2V	GBDT-XGBOOST	3	100	0.7913

Conclusion: GBDT using XGBOOST works well as compared to Random forest as XGBOOST has higher AUC value.